

1

2

3

4

5
6
7
8
9
10
11
12
13
14
15
16
17
18
19

20

Nock for Everyday Coders

Tim Galebach ~timluc-miptev
Hyperware

Abstract

Nock is not super complex, and most programmers can learn the basics of it rapidly. The mental model gained by learning Nock turns out to be very useful in learning Hoon and understanding Urbit. While the Urbit docs generally suggest to not worry about Nock, Nock is very simple and small. Most programmers will feel more comfortable in Hoon having learned Nock’s basics. The goal of this tutorial is to explain Nock clearly in terms most programmers will relate to, to impart a feeling of confidence with very basic Nock, and to give a knowledge of Nock’s idioms and big wins so that they carry over to learning Hoon.

This tutorial was originally published online by ~timluc-miptev in early 2020 and remains an excellent exposition of Nock’s affordances.

Contents

21	1 Getting Started	2
22	1.1 Phases of Nock	3
23	1.2 What Is a Nock Interpreter?	3
24	1.3 How to Run Nock Code	4
25	1.4 Subject, Formula?	4
26	1.5 Evaluating Our First Nock Code	4

27	2 Fundamental Opcodes	5
28	2.1 Nock's Simplest Functions, 0 and 1	5
29	2.2 Binary Tree Addressing	5
30	2.3 0, the "Memory Slot" Function	7
31	2.4 1, the "Quoter" Function	9
32	2.5 4, the Incrementing Function	10
33	2.6 The Cell-Maker (AKA the Distribution Rule) . .	13
34	2.7 3, the Cell Detector, and 5, the Equality Tester	15
35	2.8 2, the "Subject-Altering" Function	19
36	2.9 Summary of Fundamental Opcodes	22
37	3 Composite Opcodes	23
38	3.1 An Aside About the 1 (Quoter Function) . . .	23
39	3.2 6, "If/Else" Conditional Branch	24
40	3.3 7, the "Composition" Opcode	28
41	3.4 8, the "Variable Push" Opcode	28
42	3.5 9, Run a Stored Procedure Arm in a Core . . .	30
43	3.6 10, Replace a Memory Slot	32
44	3.7 Real Nock Code	33
45	3.8 Summary	35
46	4 Interlude	35
47	4.1 Order of Operations	35
48	4.2 11, Hints and Side Effects for the Interpreter .	36
49	4.3 Nock in Hoon	38
50	5 The Core as Design Pattern	38
51	5.1 Building a Core Manually	38
52	6 Conclusion	45

1 Getting Started

When people first look at Nock, they see the definition, which is fairly intimidating. I'm talking about lines like this:

```

56 /[(a + a) b]          /[2 /[a b]]

```

The problem is, the programmer may already know that Nock code looks more like the below – just lists of numbers, with no symbols:

```
[6
 [5 [0 6] [1 0]]
 [0 7]
 [9 4 [[0 2] [2 [0 6] [0 5]] [4 0 7]]]
]
```

What gives? Which one is the “real” Nock?

1.1 Phases of Nock

We are looking at two different things in the examples above:

1. Pseudocode for *how to interpret* Nock.
2. Code to be interpreted (written as lists of numbers).

The symbols and spec are pseudocode, not real Nock code. They could just as easily be written in English, and they will never be written down as actual Nock code and given to an interpreter. They represent what an interpreter should do to turn Nock code into interpreter instructions.

The lists of numbers are the actual Nock code. This is what you feed to an interpreter to get some result.

1.2 What Is a Nock Interpreter?

An interpreter can be a computer program, or it can be a human manually expanding Nock code into results. In both cases, the program and human have to know the Nock pseudocode in order to do the right thing with incoming Nock code.

So a Nock interpreter is any entity that takes Nock code as input, and gives a noun as output. A noun can be:

```
:: a number
782
:: a cell (pair with two elements)
[782 9872]
:: each element can itself be a pair
```

```

89 [782 [9872 89728]]
90 :: the above can be written as
91 [782 9872 89728]

```

92 1.3 How to Run Nock Code

93 We will be expanding Nock pseudocode manually in the exam-
 94 ples that follow, in effect acting as our own interpreter.

95 If we want to check that we’re getting the right results from
 96 our manual interpretation, we need to run a Nock interpreter,
 97 such as the Urbit Dojo.

- 98 1. Start up a Dojo session on a fake ship.¹
- 99 2. At the prompt, we can execute Nock using `. * dottar`, e.g.,
 100 `. * (NOCK_SUBJECT, NOCK_FORMULA).`

101 1.4 Subject, Formula?

102 Let’s keep this simple:

- 103 • subject = an argument to a function
- 104 • formula = the function

105 That’s it. We’ll see below how this works, going really
 106 slowly with examples.

107 1.5 Evaluating Our First Nock Code

108 OK, so the interpreter takes two arguments, a “subject” and a
 109 “formula”. Both are nouns (a number or a cell). Let’s run some
 110 insanely simple Nock code in the Dojo:²

```

111 > . * (42 [0 1])
112 42

```

¹Consult <https://docs.urbit.org> for details.

²Code samples beginning with `>` are Dojo inputs (to wit, Hoon expressions).

In the above, 42 is our subject. [0 1] is our formula.

Formulas are always cells, and the first element of the cell is a number that you can think of as *the name of the function*.

In this case, our function name is 0, which is the memory slot function. It is always followed by 1 number, in this case 1, which is the number of the memory slot to fetch in the subject.

Whenever we look at Nock code, we want to ask:

- What is the subject (function argument)? In this case, it's 42.
- What is the formula (function)? In this case, it's [0 1].
- What value does that formula (function) produce when called on this subject (argument)? In this case, the return value is 42.

Why is the return value 42? How does this formula work?

2 Fundamental Opcodes

2.1 Nock's Simplest Functions, 0 and 1

The two most basic Nock functions are 0 address and 1 constant. The goal here is to get strong intuitions of what they do, how they handle edge cases, and how this relates to the Nock spec/pseudocode.

2.2 Binary Tree Addressing

Before getting started on Nock proper, we should understand how Nock and Hoon handle addresses in binary trees. If you already understand why memory slot 5 of [['apple' %pie] [0b1101 0xdad]] is %pie, you are good to go and can skip ahead to the Section 2.3.

Every noun in Nock can be thought of as a tree, which means we can give an exact number to access any position in the tree. This means that, no matter how big our subject (argument) is, we can yank a value out of any part of it.

The Nock specification defines noun tree addressing:


```

144
145 / [1 a]          a
146 / [2 a b]       a
147 / [3 a b]       b
148 / [(a + a) b]   / [2 / [a b]]
149 / [(a + a + 1) b] / [3 / [a b]]
150

```

151 This permits the address to be defined within a particular sub-
152 tree as well as within the overall tree (Figure 1).

Figure 1: A binary tree with labeled node addresses to several layers.



153 That is, how do you say which slot number you want from
154 a given tree? We say that:

- 155 • The tree root is address 1.
- 156 • The head of every node n is $2n$.
- 157 • The tail of every node n is $2n + 1$.


158 Let's take an example tree to illustrate. In Nock cell form,
159 the tree is:

160 `[[4 5] [6 14 15]]`

161 Diagrammatically, the tree looks like Figure 2.

- 162 • 1 is the address of the whole tree, `[[4 5] [6 14 15]]`.
- 163 • 2 is the address of the left branch, `[4 5]`.

Figure 2: A binary tree with some labeled node addresses.



- 3 is the address of the right branch, [6 14 15].
- 15 is the value 15.

Let's play around now in the Dojo Nock interpreter so that we can confirm this. In each example, our subject (argument) will be the tree [[4 5] [6 14 15]].

```

:: formula (function) [0 1]: get the whole tree
> .*([[4 5] [6 14 15]] [0 1])
[[4 5] 6 14 15]

:: formula (function) [0 2]: get the left branch
> .*([[4 5] [6 14 15]] [0 2])
[4 5]

:: formula (function) [0 7]: get the subtree in slot 7
> .*([[4 5] [6 14 15]] [0 7])
[14 15]
```

2.3 0, the "Memory Slot" Function

The pseudocode for the 0 opcode³ is as follows:

```
*[a 0 b] / [b a]
```

³From this point on, we will write Nock opcodes in special type so that they can be distinguished from other numbers in Nock. This reflects Hoon practice, which marks Nock rules as constants like %0. While not strictly necessary, it can be helpful when learning or coding Nock to see an explicit distinction.

187 /[b a] is pseudocode. In English, it means “a is a binary
188 tree. Get the memory slot numbered b.

189 So if a were the tree [9 10], and b were 1, we’d get the
190 memory slot 1 in the tree [9 10], which is just the tree itself.

191 Written in pseudocode, that’s /[1 [9 10]]. We can also
192 do /[2 [9 10]], which grabs memory slot 2, aka 9.

193 Let’s look at some examples. I’ve put them first in Dojo
194 form so that you can see how they run in that interpreter, and
195 then I show the “human interpreter” in pseudocode below that.

196 2.3.1 Example: Retrieve the Subject

```
197
198 :: get memory slot 1
199 > .*([50 51] [_ 1])
200 [50 51]
201 :: PSEUDOCODE -- replaces the Dojo's () with []
202 *[[50 51] [_ 1]]
203 :: *[a @0@ b] -> a = [50 51], b = 1
204 /[1 [50 51]]
205 [50 51]
```

207 2.3.2 Example: Retrieve the Head of the Subject

```
208
209 > .*([50 51] [_ 2])
210 50
211 :: PSEUDOCODE -- replaces the Dojo's () with []
212 *[[50 51] [_ 2]]
213 :: *[a @0@ b] -> a = [50 51], b = 2
214 /[2 [50 51]]
215 50
```

217 2.3.3 Example: A Crash

```
218
219 > .*([50 51] [_ [0 1]])
220 dojo: hoon expression failed
221 :: PSEUDOCODE
222 *[[50 51] [_ [0 1]]]
223 /[[0 1] [50 51]]
```



```

224 :: can't evaluate this--a memory slot must be a
225     number like 2, not a cell like [0 1]
226 CRASH
227

```

2.3.4 Summary of @

@ is everywhere in Nock, because “get something from a memory slot” really means “store stuff in a place and get it whenever I want,” which is another name for creating variables. These memory slot numbers take the place of variable names.

If you’re familiar with assembly or C, they’re conceptually similar to memory pointers or registers in how Nock uses them. Keep in mind though, Nock is functional/immutable, so it doesn’t update memory locations: it creates copies of data structures with altered values.

We’ve also seen that the memory slot function can’t take just anything as the memory slot to fetch: it must receive an atom (number).

2.4 1, the “Quoter” Function

This is another really simple function. You can think of it as a “quoter”: it just returns any value passed to it exactly as it is. It ignores the subject, and just quotes the value after it. Let’s look at a couple examples.

```

246 > .*([20 30] [1 67])
247 67
248 :: [@1@ 2 587] is same as [@1@ [2 [587]]]
249 > .*([20 30] [1 [2 [587]]]
250 [2 587])

```

It doesn’t matter how much information is after the 1: 1 is a dumb function that just returns it all.

The pseudocode for 1 is:

```

254 * [a 1 b]      b
255
256

```

In English, this means: “ignore the subject ‘a’, and just return everything after the 1 exactly as it is.

Let's look at our first example, `.*([20 30] [1 67])`. The subject `a` is `[20 30]`, so we ignore that. What's after the `1`? `67`, so we return that.

In the second example, "everything after the `1`" is longer, but the same rule applies: the Nock interpreter just returns it exactly as it is, after stripping out unnecessary brackets.

2.4.1 Summary of `1`

`1` is a simple function that just returns whatever is after it ("quoter" or "constant"). It's useful for quoting values that you want to use later in your Nock code.

2.5 `4`, the Incrementing Function

`0` and `1` are simple functions that don't have any nested behavior. Now we're going to move to a function that *does* have nested children, opcode `4`. In these examples, pay attention to how `4` operates on its arguments; we'll look at pseudocode and break down the examples in a moment.

```
> .*(50 [4 0 1])
51
> .*(50 [4 4 0 1])
52
> .*([100 150] [4 4 0 3])
152
> .*(50 [4 1 98])
99
> .*(50 [4 1 [0 2]])
dojo: hoon expression failed
```

Here's `4`'s pseudocode, juxtaposed with that of `0` and `1`:

```
*[a 4 b]  +*[a b]
*[a 0 b]  /[b a]
*[a 1 b]  b
```

295 In English, this says “when we have subject a, function 4,
 296 and Nock code b, first evaluate [a b] as [subject formula],
 297 and then add 1 to the result.”

298 If we contrast with 0 and 1, we see that the right side has
 299 a * symbol. This symbol means “evaluate the expression again
 300 in the Nock interpreter.” 0 and 1 *did not* have this symbol, and
 301 that’s why they couldn’t evaluate nested Nock expressions.

302 2.5.1 Example: Walking Through Increment

303 Let’s start by translating the Dojo’s `.*(subject formula)` to
 304 pseudocode of the form `*[a b]`, and then expand it line by
 305 line:

```
306 > .*(50 [4 0 1])
307 51
308 :: change to pseudocode
309 *[50 [4 0 1]]
310 :: move the 4 to the outside as +
311 **[50 [0 1]]
312 :: expand the 0 opcode to the memory slot operator "/"
313 +/[1 50]
314 :: grabs the memory slot
315 +(50)
316 :: evaluate
317 51
```

318 The [a b] part of `**[a b]` expands to:

319 [50 [0 1]]

320 OK, this we know how to handle! It’s just our 0 function, and
 321 it wants the value in memory slot 1 of the subject. That’s 50.

322 Now we know that `*[a b]`=50, and we just have to add 1
 323 to it (the + in `**[a b]`). That is 51, which is exactly what the
 324 interpreter gave us.

325 2.5.2 Example: Walking Through Increment

326 This one is similar, we just have an extra 4. We again start by
 327 translating the Dojo’s `.*(subject formula)` to pseudocode,
 328 and then expand

```

329 > .* (50 [4 4 0 1])
330 52
331 *[50 [4 4 0 1]]
332 :: the first 4 moves outside as a '+'
333 **[50 [4 0 1]]
334 :: 2nd 4 becomes a '+'...we have opcode 0 again!
335 ***[50 [0 1]]
336 ++/[1 50]
337 ++(50)
338 +(51)
339 52

```

340 2.5.3 Example: Walking Through Serial Increment

341 Here we again see lots of 4s applied consecutively, and we also
342 see how we can yank values out of a more complicated subject
343 and manipulate them. Notice how the subject is a cell, not an
344 atom. The rest is the same as in the previous example.

```

345 > .* ([100 150] [4 4 0 3])
346 152
347 a (the subject) = [100 150]
348 *[[100 150] [4 4 0 3]]
349 **[[100 150] [4 0 3]]
350 :: Now we've extracted all the increments,
351 :: so we just grab the value at memory slot 3
352 ***[[100 150] [0 3]]
353 ++/[3 [100 150]]
354 ++(150)
355 +(151)
356 152

```

357 2.5.4 Example: Walking Through Incrementing a Constant

358 In the below example, we see how we can use the quote/con-
359 stant function 1 to generate the value 98 and increment it. We
360 ignore the subject 50 completely.

```

361 > .* (50 [4 1 98])
362 99
363 *[50 [4 1 98]]
364 :: formula is the "quoter" function

```

```

365  ++[50 [1 98]]
366  +(98)
367  99

```

368 2.5.5 Example: A Crash When Incrementing a Cell

369 Just as opcode 0 had values it couldn't handle (non-atoms), so
 370 opcode 4 needs the nested value inside it to evaluate to an atom.

```

371  > .* (50 [4 1 [0 2]])
372  dojo: hoon expression failed
373  *[50 [4 1 [0 2]]]
374  :: OK cool, the nested value is a 1 opcode
375  ++[50 [1 [0 2]]]
376  :: 1 ignores the subject (50) and just returns [0 2]
377  +([0 2])
378  :: [0 2] isn't an atom, so how can we increment it???
379  :: We can't, so we crash.
380  dojo: hoon expression failed

```

381 2.5.6 Summary of 4

382 In these examples, we've seen that function 4 can be called as
 383 many times in a row as we want. At the end of those calls, it
 384 always ends up incrementing a number that either is yanked
 385 from the subject (memory slot function 0) or quoted as it is
 386 (quote function 1).

387 2.6 The Cell-Maker (aka the Distribution Rule)

388 The Nock interpreter is allowed to return nouns, which are
 389 atoms (positive numbers) or cells (pairs of nouns). What have
 390 our functions/opcodes been returning so far?

- 391 • 0: atoms or cells, depending what's in the memory slot
- 392 that we yoink.
- 393 • 1: atoms or cells, depending on what we quote.
- 394 • 4: just atoms.

But what if my subject was [51 67 89], and I wanted to increment every value and return that as [52 68 90]? How can I do that when it's a cell, and 4 only seems able to return atoms?

The answer is something that the Nock docs call the “distribution rule” or “implicit cons” (hello, fellow LISPer!), but that I find easiest to think of as the “Cell-Maker Rule”.

2.6.1 A Quick Detour into Nock Formulas

We haven't talked much yet about what values are legal to feed into the Nock interpreter (the `.*(subject formula)` function in the Dojo). So far, we've only been using formulas that start with numbers (our functions/opcodes 0/1/4).

Let's now fully solidify our understanding of what's a legal formula by taking a quick look at the three possible cases. (The format is `*(subject formula)`.)

- The formula is a cell starting with an opcode. We know this is OK.

```
> .*(50 [0 1])
50
```

- The formula is just an atom (o). This is not valid.

```
> .*(50 0)
dojo: hoon expression failed
```

- Finally, we try a formula that is a cell starting with an atom. This looks invalid, but—it works!

```
> .*(50 [[0 1] [1 203]])
[50 203]
```

So apparently a formula cell can start with a cell. The Cell-Maker rule is:

```
*[subject [formula-x formula-y]]→
  [*[subject formula-x] *[subject formula-y]]
```

In our example above, `formula-x` is `[0 1]`, and `formula-y` is `[1 203]`. They each evaluate individually against the subject, and the end result is a cell.

We can make as many cells in a row as we want:

```
> .* (50 [[0 1] [1 203] [0 1] [1 19] [1 76]])
[50 203 50 19 76]
```

We can put any operation inside each cell:

```
> .* ([19 20] [[0 1] [1 76] [4 4 0 3]])
[[19 20] 76 22]
```

If we take the returned collection `[[19 20] 76 22]` in order, we can write in English how they connect to our collection of formulas that we passed:

- `[0 1]`: get memory slot 1
- `[1 76]`: return the quoted value 76
- `[4 4 0 3]`: increment twice the value in memory slot 3 (20)

So we can pass one small subject `([19 20])` and make an arbitrarily long collection of values from it, using any functions we want. Cell-Maker FTW!

2.7 3, the Cell Detector, and 5, the Equality Tester

Now we come to functions/opcodes 3 and 5, which are pretty straightforward after we've seen how 4 and the Cell-Maker work. Functions 3 and 5, like 4, allow nested evaluation. Let's put all their pseudocode definitions together to compare:

```
:: function/opcode 3
*[a 3 b]    ?*[a b]
:: function/opcode 4
*[a 4 b]    +*[a b]
:: function/opcode 5
*[a 5 b c]  =*[a b] *[a c]
```

459 First of all, notice how the right side of all these “equations”
 460 has the evaluation operator, `*`. This means that these functions
 461 can have nested formulas, since they keep evaluating all the
 462 way down.

463 There are some new pseudocode symbols here that we need
 464 to translate into English. We already know `+`: “increment the
 465 value after this”. Now we also see:

- 466 • `?`: “check whether the value after this is a cell. Return 0
 467 if yes, 1 if no”.⁴
- 468 • `=`: “first run the function in `b` with subject `a` as the argu-
 469 ment, and same for the function in `c`. If the results are
 470 equal, return 0, if not, return 1.

471 2.7.1 Example: Not a Cell

```
472 > .* (50 [3 0 1])
473 1
474 :: PSEUDOCODE OF THE ABOVE
475 *[50 [3 0 1]]
476 ?*[50 [0 1]]
477 ::get memory slot 1 of the subject
478 ?(50)
479 :: is 50 a cell? No, so return 1
480 1
481
482
```

483 2.7.2 Example: A Cell

```
484 > .* ([50 51] [3 0 1])
485 0
486 :: PSEUDOCODE OF THE ABOVE
487 *[[50 51] [3 0 1]]
488 ?*[[50 51] [0 1]]
489 ::get memory slot 1 of the subject: [50 51]
490 ?([50 51])
491 :: is [50 51] a cell? Yes, so return 0
492 0
493
494
```

⁴Nock conventionally explains that there is one “true” case and (potentially) many “false” cases, thus `0/true` and `1/false`.

2.7.3 Example: Nested Cell Evaluation with 4

```

495
496 > .*([50 51] [4 4 3 0 1])
497 2
498
499 :: PSEUDOCODE OF THE ABOVE
500 *[[50 51] [4 4 3 0 1]]
501 **[[50 51] [4 3 0 1]]
502 :: whatever comes out of the 3 function,
503 :: we're gonna increment it twice
504 ***[[50 51] [3 0 1]]
505 :: down to just fetching memory slot 1
506 ++*[[50 51] [0 1]]
507 ++?([50 51])
508 :: is [50 51] a cell? Yes, so return 0
509 ++(0)
510 +(1)
511 2
512

```

2.7.4 Example: Check Multiple Cases of Cells

```

514
515 > .*([[50 51] 52] [[3 0 2] [3 0 3]])
516 [0 1]
517 *[[50 51] 52] [[3 0 2] [3 0 3]]
518 ?[*[[50 51] 52] [0 2]]
519    *[[50 51] [0 3]]
520 :: yank memory slots 2 and 3
521 ?*[[50 51] 52]
522 :: first is a cell, second is not
523 [0 1]
524

```

2.7.5 Example: Equality Test

Because 5 compares the results of 2 formulas, it *always* makes 2 inner evaluations of the subject. It's similar to Cell-Maker in this way.

```

529
530 > .*([50 51] [5 [0 2] [0 2]])
531 0
532 :: PSEUDOCODE
533 *[[50 51] [5 [0 2] [0 2]]]

```

```
534 :: factor out the =
535 =[*[[50 51] [0 2]] *[[50 51] [0 2]]]
536 :: get memory slot 2 twice
537 =(50 50)
538 0
```

540 2.76 Example: Comparing Two Unequal Values

```
541
542 > .*([50 51] [5 [0 2] [0 3]])
543 1
544 :: PSEUDOCODE
545 *[[50 51] [5 [0 2] [0 3]]]
546 :: factor out the =
547 =[*[[50 51] [0 2]] *[[50 51] [0 3]]]
548 :: get memory slot 2 and memory slot 3
549 =(50 51)
550 1
551
```

552 2.77 Example: Compare Values with Function Calls

```
553
554 > .*([50 51] [5 [4 0 2] [0 3]])
555 0
556 :: PSEUDOCODE
557 *[[50 51] [5 [4 0 2] [0 3]]]
558 :: factor out the =
559 =[*[[50 51] [4 0 2]]
560   *[[50 51] [0 3]]]
561 :: factor out the +
562 =[+*[[50 51] [0 2]]
563   *[[50 51] [0 3]]]
564 :: get memory slots 2 and 3
565 =[+50 51]
566 :: evaluate the +
567 =(51 51)
568 0
569
```

570 2.78 Example: Compare Cells with Function Calls

```

571
572 > . * ([99 99] [5 [1 [99 99]] [0 1]])
573 0
574 :: PSEUDOCODE
575 * [[99 99] [5 [1 [99 99]] [0 1]]]
576 = [* [[99 99] [1 [99 99]]]
577    * [[99 99] [0 1]]]
578 :: first cell is the result of quoter function
579 :: second cell is the result of memory fetch
580 = [[99 99] [99 99]]
581 :: true
582 0
583

```

584 2.7.9 Summary of 3 and 5

585 These along with 4 are known as the “axiomatic” functions.
 586 They exist to implement definitional logic for the Nock spec-
 587 ification. In particular, the cell check and the equality check
 588 provide for structural analysis of nouns.

589 2.8 2, the “Subject-Altering” Function

590 In all our examples so far, the subject has been defined at the
 591 start when we call the interpreter, and never changes. But what
 592 if we want to change the subject?

593 A different subject? Why would we want that? Here’s an
 594 easy example. Say you found the following piece of Nock code
 595 on the interwebz:

```

596 [8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2]
597    [4 0 6] 0 7] 9 2 0 1]

```

598 This is the code for a Nock function that expects a subject
 599 that is an atom, and decrements that subject by 1. You can
 600 actually enter it in the Dojo right now:

```

601 > . * (100 [8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0
602    2] [4 0 6] 0 7] 9 2 0 1])
603 99

```

604 This works! You do not have to understand the code right now;
 605 just try entering different numbers instead of 100 to see the
 606 program working.

607 But now imagine that I have a Nock program with a differ-
608 ent subject

609 > .*([50 51] some-formula)

610 And somewhere inside `some-formula`, I want to decre-
611 ment the number 51. I can't pass `[50 51]` as the subject to
612 my decrementer code above though, because that's a cell, and
613 it expects just an atom (a number).

614 2.8.1 Subject Altering to the Rescue

615 The function/opcode `2` is designed to handle this problem for
616 us. First the pseudocode:

617 `*[a 2 b c] *[*[a b] *[a c]]`
618
619

620 `2` expects 2 formulas after the subject `a`: `b` and `c`. With
621 those, it:

- 622 1. runs formula `b` against the subject to set up a new envi-
623 environment/subject derived from the subject
- 624 2. runs formula `c` against the subject to prepare a 2nd func-
625 tion
- 626 3. run that 2nd function against the new environment/sub-
627 ject from step (1)

628 Note that the pseudocode for `2` has nested “`*`”s.

629 `*[*[a b] *[a c]]`

630 The two inner `*s` run steps (1) and (2), and the outer one,
631 around the whole expression, runs step (3).

632 2.8.2 Example: Change the Subject and Call a Constant Value

633 > .*([50 51] [2 [0 3] [1 [4 0 1]]])
634
635 52
636 :: PSEUDOCODE
637 `*[[50 51] [2 [0 3] [1 [4 0 1]]]`
638 `:: separate b and c to each run against the subject`
639 `(steps 1 and 2)`

```

640 *[*[[50 51] [0 3]] *[[50 51] [1 [4 0 1]]]]
641 :: after steps 1 and 2, we have a new subject, 51!
642 :: note how we're back in normal *[subject formula]
643   form
644 *[[51 [4 0 1]]
645 :: apply the 4 function as we're used to
646 **[[51 [0 1]]
647 :: grab 51 from memory slot 1
648 +(51)
649 52
650

```

651 2.8.3 Example: Grab a Block of Code from the Subject and Run It

652 Think of this as grabbing a “stored procedure” from the subject.

```

653
654 > .*([[4 0 1] 51] [2 [0 3] [0 2]]])
655 52
656 :: PSEUDOCODE, subject is [[4 0 1] 51]
657 *[[[4 0 1] 51] [2 [0 3] [0 2]]]
658 *[*[[[4 0 1] 51] [0 3]
659   *[[[4 0 1] 51] [0 2]]]]
660 :: step 1 gets memory slot 3, step 2 grabs memory
661   slot 2
662 *[[51 [4 0 1]]
663 :: looks like a normal 4 opcode to me!
664 **[[51 [0 1]]
665 :: grab memory slot 1
666 +(51)
667 52
668

```

669 2.8.4 Example: Back to Our Motivating Case

670 Remember our decrementing block of code that we couldn't
671 use when the subject was [50 51], instead of just an atom?
672 Opcode 2 makes handling that issue a piece of cake.

673 We simply use 2 to transform our subject into an atom, and
674 use 1 to quote the decrement block of code before it evaluates
675 in step (3).

```

676
677 > .*([50 51] [2 [0 2] [1 [8 [1 0] 8 [1 6 [5 [0 7] 4 0
678   6] [0 6] 9 2 [0 2] [4 0 6] 0 7] 9 2 0 1]]])

```

```

679 49
680 :: PSEUDOCODE (subject is [50 51])
681 :: ("decrement-formula" substituted for clarity)
682 *[[50 51] [2 [0 2] [1 decrement-formula]]]
683 *[*[[50 51] [0 2]] *[[50 51] [1 decrement-formula]]]
684 *[[50 decrement-formula]
685 ::decrement-formula has the atom subject that it wants
686 49
687

```

688 2.8.5 Summary of 2

689 For those who know a bit of Hoon, the second example above
690 is rather similar to calling an arm that produces a gate, and
691 then running the gate. Most of Hoon runs this type of stored-
692 procedure plus subject-altering Nock, and it all uses opcode 2
693 at its base.

694 And for those who know Hoon, `[[4 0 1] 51]` should al-
695 ready be looking a lot like `[battery payload]...` that's not
696 a coincidence. We're starting to see the first glimpses of how
697 Hoon's cores, arms and subjects/subject mutations flow out of
698 the fundamental structure of Nock.

699 We also see how Hoon/Nock lend themselves well to throw-
700 ing around chunks of code, and adjusting the subject as neces-
701 sary to create the correct subject/environment against which
702 to run that code.

703 2.9 Summary of Fundamental Opcodes

704 You now have seen all the fundamental functions/opcodes in
705 Nock. In the next part, I'll introduce the remaining functions,
706 which no longer need pseudocode: we have enough scaffold-
707 ing now to build the rest of Nock from Nock itself. Instead,
708 these new opcodes will just be shortcuts/macros/code expan-
709 sions building on opcodes 0-5.

710 We'll also start to connect Nock to Hoon, and see how most
711 of the fundamental (and slightly weird) features of Hoon flow
712 directly from Nock's structure, and make a lot more sense in
713 combination with it.

714 3 Composite Opcodes

715 A word of encouragement: we're done with the hard part now.
716 Every Nock function we learn in this section will be built from
717 pieces in the preceding section.

718 None of these new functions are *necessary* to make Nock
719 work. All of them except Nock 10 and 11 are syntactic sugar
720 that is part of the Nock definition, and must be built into every
721 correct Nock implementation. If you have seen macros or code
722 expansions in other languages, that's another word for what's
723 happening here.

724 Nock 10 makes it easier to replace a memory value some-
725 where in a tree, and Nock 11 allows passing hints to the inter-
726 preter.

727 One point of clarity: this syntactic sugar/code-expansion
728 system is *not* extensible. This means that you can't invent your
729 own Nock opcodes and still have that language be Nock; you're
730 making a higher-level language on top of Nock at that point.

731 In fact, that's not a bad way to think of Hoon: it's a higher-
732 level language that adds missing syntactic sugar/macros and
733 human-readable names to Nock. (That's not the whole story,
734 but it's a decent mental peg to initially hang Hoon on.)

735 Nock is intentionally very, very small, such that you can
736 always walk through and analyze what is happening in a block
737 of code if you know Nock's opcodes.

738 3.1 An Aside About the 1 (Quoter Function)

739 You're going to see the 1 opcode (the "Quoter") appear in a lot
740 of examples below for a simple reason: the 6-9 opcodes expect
741 formulas in a lot of places, not just atoms. And, as we've seen
742 already, formulas have to be cells.

743 Whenever a formula is required, but you really just want
744 to return a number, you use the quoter function.

745 3.1.1 Example: Incrementing a Constant

746 :: We just want to run 4 on the number 5,
747

```

748 :: but 4 expects a formula after it, so we use [1 5]
749 > .* (0 [4 1 5])
750 6
751

```

752 3.1.2 Example: Comparing a Memory Slot to a Constant Value

```

753
754 :: we grab memory slot 2
755 :: then it has to be compared to the result of a
756     formula
757 :: so we just use the formula [1 23] to return 23
758 > .* ([23 45] [5 [0 2] [1 23]])
759 0
760

```

761 3.2 6, "If/Else" Conditional Branch

762 The remaining Nock opcodes are “sugar”. This means that the
763 functions in this next part will use only `*` and Nock code in
764 their pseudocode. For example, here is the code for 6, the
765 “If/Else” function.

```

766
767 * [a 6 b c d]      * [a * [[c d] 0 * [[2 3] 0 * [a 4 4 b]]]]
768

```

769 The prose explanation of what’s happening is straightforward:
770

- 771 1. Evaluate `b` against the subject `a` (`* [a b]`) to see whether
772 it’s `0/true` or `1/false`.
- 773 2. If `b` equals `0/true`, run formula `c` against subject `a`.
- 774 3. If `b` equals `1/false`, run formula `d` against subject `a`.
- 775 4. If `b` is not equal to `0` or `1`, the code crashes, for reasons
776 we’ll see in the code explanation.

777 3.2.1 Code Explanation (Way More Fun)

778 OK, so that’s the English version. The code explanation (the
779 right side of the definition) is *really fun* now that we know the
780 basic Nock opcodes.

781 The pseudocode has four nested “[subject formula]’s, so
 782 I’m going to unwrap those to the bottom, and then build it up
 783 again. The layers are, in order:

- 784 1. `*[a ...]`, i.e. subject `a` evaluated against that long for-
 785 mula starting with `*[c d] 0 ...]`.
- 786 2. `*[c d]`, i.e. subject `c` evaluated against the formula start-
 787 ing with `[2 3 ...]`.
- 788 3. `*[2 3]`, i.e. subject `2` evaluated against the lowest-level
 789 formula.
- 790 4. Finally, subject `a` evaluated against formula `[4 4 b]`.

791 3.2.2 Step 4

792 Remember, our English explanation was “see if `*[a b]` is true
 793 or false, and do different actions depending on that. Step Four
 794 is that check.

795 Let’s say we have `a` as `59`, and `b` as the quoted value `0/true`.

```
796 *[a 4 4 b]
797 :: a: 59
798 :: b: [1 0]
799 *[59 4 4 [1 0]]
800 :: substitute out the two increment operators
801 ++*[59 [1 0]]
802 :: ignore subject, return quoted value 0
803 ++(0)
804 2
```

805 In summary, because `*[a b]` evaluated to `0/true`, we get
 806 the number `2`. If `*[a b]` had been `1/false`, we’d get `3` (because
 807 we’d evaluate `++(1)`).

808 What the heck? Why are we getting `2` or `3` back? How
 809 does that help us?? Well, Nock uses that returned `2` or `3` as a
 810 *memory slot*, and executes the code in that memory slot.

811 3.2.3 Step 3

812 Now we go up to step 3, with the subject `[2 3]` evaluated
 813 against our return value from step 4. Let’s imagine that `2` had

814 been returned:

```
815 :: remember, "result-of-step-4" was 2
816 *[[2 3] @ result-of-step-4]
817 *[[2 3] @ 2]
818 :: get memory slot 2
819 2
```

820 This grabs memory slot 2 if $[a \ b]$ was true, and memory
821 slot 3 if $[a \ b]$ was false.

822 Wait, isn't this redundant? We are just using our 2 or 3
823 generated in step 4 to generate a 2 or a 3. Seems dumb.

824 The answer is that we make sure that a crash happens if
825 $[a \ b]$ yields any answer other than 0/true or 1/false. If $[a \ b]$
826 returned 10, for example, we'd have the following code in
827 step 3:

```
828 *[[2 3] @ 10]
829 :: there's no slot 10
830 CRASH!!
```

831 This is exactly what we want: the program crashes unless
832 we are doing a loobean⁵ test that returns a 0 or 1 (converted to
833 an indicial 2 or 3) in step 4.

834 3.24 Step 2

835 We now have our validated 2 or 3 to plug into step 2. Let's
836 imagine c is the simple formula $[0 \ 1]$ and d is $[1 \ 203]$.

```
837 :: if step 3 returned "2" (true)
838 *[[[0 1] [1 203]] @ 2]
839 [0 1]
```

840 Not much to see here: we just grab memory slot 2 or 3
841 depending on whether our initial b was true or false.

842 3.25 Step 1

843 And now we're back at the top level, where we just use whichever
844 formula we yonked in step 2 and run it against a .

⁵A boolean truth value but with the typical truth values of 0 and 1 reversed to 0/true and 1/false.

```

845 *[a formula-from-step-2]
846 :: let's say we returned formula [0 1]
847 :: our original a, from step 4, was 59
848 *[59 [0 1]]
849 59

```

850 3.2.6 Example: Code Expansion of 6

```

851 > .*(1 [6 [0 1] [0 1] [4 0 1]])
852
853 :: PSEUDOCODE
854 :: *[a *[[c d] 0 *[[2 3] 0 *[a 4 4 b]]]]
855 *[*1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 *[*1 4 4 [0 1]]]]]
856 :: factor out the 4 opcodes
857 *[*1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 +*[*1 [0 1]]]]]
858 :: b evaluates to 1 (yank memory slot 1)
859 *[*1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 ++(1)]]]
860 :: evaluate the two increments
861 *[*1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 3]]]
862 :: get memory slot 3 from [2 3]
863 *[*1 *[[[0 1] [4 0 1]] 0 3]]
864 :: [0 3] means get memory slot 3 from the subject
865      (formula [4 0 1])
866 *[*1 [4 0 1]]
867 :: factor out the 4 opcode
868 +*[*1 [0 1]]
869 +(1)
870 2
871

```

872 3.2.7 Summary of 6

873 Now I'm going to make you a little sad. Most Nock interpreters
874 don't do this whole awesome code expansion. They just see 6,
875 and implement an if/else check with a crash if `*[a b]` isn't a
876 loobean.

877 However, the pattern of storing chunks of code in memory
878 and pulling them out when you want them is the most impor-
879 tant part of Nock. In fact, those of you who know Hoon are
880 probably already seeing the makings of the “core” code pat-
881 tern.

3.3 7, the "Composition" Opcode

7 is so simple we barely need to spend time on it. All it does is create a new subject/environment (using 2), and immediately runs a formula against that new subject. Let's compare it to 2:

```

:: opcode 7
*[a 7 b c]  *[*[a b] c]

:: opcode 2
*[a 2 b c]  *[*[a b] *[a c]]

```

This is almost exactly the same except with c instead of *[a c]. This means that you can write a "subject-changing" formula in b, and then run a simple function against that in c.

3.3.1 Example: Opcode 2 vs. Opcode 7

```

:: using opcode 2
> .*([23 45] [2 [0 3] [1 4 0 1]])
46

:: using opcode 7 -- we get to remove the quoter
  opcode "1"
:: wow amazing /sarcasm
> .*([23 45] [7 [0 3] [4 0 1]])
46

```

Opcode 7 is clearly trivial, so why does it exist? It allows clean expression of function composition: this is really $c(b(a))$, where a (the subject) is the initial argument, and b and c are functions. This pattern comes up a ton, and it's nice to not use 1s everywhere.

3.4 8, the "Variable Push" Opcode

If you're ever writing Nock and think "how can I add a new variable to the subject?", opcode 8 is what you want. The new variable can be based on either the existing subject, or be a new value you add on.

```

*[a 8 b c]          *[[*[a b] a] c]

```

921 This says to run `*[a b]`, and then make that the head of a
 922 new subject, with the old subject `a` as the new subject's tail.

923 3.4.1 Example: Add Variable as a Copied Value from an Existing Sub- 924 ject

925 In English, the below code first yanks the variable from mem-
 926 ory slot 3, copies it to the head of a new subject, and then in-
 927 crements the value in that new head.

```

928 > .*([67 39] [8 [0 3] [4 0 2]])
929
930 40
931 :: PSEUDOCODE
932 *[[*([67 39] [0 3]) [67 39]] [4 0 2]]
933 :: yanks mem slot 3 and pins it to the front of the
934 old subject
935 *[[39 [67 39]] [4 0 2]]
936 **[[39 [67 39]] [0 2]]
937 +(39)
938 40
939
  
```

940 3.4.2 Example: Add Variable as a New Value

```

941 > .*([67 39] [8 [1 0] [4 0 2]])
942
943 1
944 :: PSEUDOCODE
945 *[[*([67 39] [1 0]) [67 39]] [4 0 2]]
946 *[[[0 [67 39]] [4 0 2]]
947 **[[[0 [67 39]] [0 2]]
948 +(0)
949 1
950
  
```

951 Why do we want this? In the first example, we pin a copy
 952 of a value so that we can manipulate it without changing the
 953 original. In the second example, we add a `0` to the front; maybe
 954 we want to increment it until some condition is met?

955 The above should be starting to feel *very* Hoon-ish: we're
 956 a minor code transform away from pinning new values to the
 957 head of a payload.

958 3.5 9, Run a Stored Procedure Arm in a Core

959 We're almost at full Hoon now, although still at a very low/raw
960 level. 9 looks a little complicated...

961 $\star[a \ 9 \ b \ c] \ \star[\star[a \ c] \ 2 \ [0 \ 1] \ 0 \ b]$
963

964 ... but in English, this is just saying:

- 965 1. Use formula c to make a new subject from $a \ (\star[a \ c])$.
- 966 2. Grab the formula located at memory slot b in that new
967 subject.
- 968 3. Run that formula against the new subject $\star[a \ c]$.

969 The fact that the above description has the words “make a new
970 subject” tells us right away that it’s syntactic sugar for opcode
971 2, and that’s indeed what we see in the pseudocode.

972 3.5.1 Example: Using 9 to Run an Increment Arm

973 The initial expression here likely looks cryptic, but if you fol-
974 low the pseudocode, it will become clear.

975 To stay oriented, remember that, if we’re thinking of 9 as
976 $\star[a \ 9 \ b \ c]$, then

977 $a: 45$
978 $b: 2$
979 $c: [[1 \ 4 \ 0 \ 3] \ 0 \ 1]$

```

980 > .*(45 [9 2 [1 4 0 3] 0 1])
981 46
982
983 :: PSEUDOCODE
984  $\star[45 [9 \ 2 \ [1 \ 4 \ 0 \ 3] \ [0 \ 1]]]$ 
985  $\star[\star[45 [1 \ 4 \ 0 \ 3] \ [0 \ 1]] \ 2 \ [0 \ 1] \ 0 \ 2]$ 
986 :: new subject is  $[[4 \ 0 \ 3] \ 45]$ 
987 :: that is a formula to increment mem slot 3 in the
988 :: head; 45 in the tail
989  $\star[[[4 \ 0 \ 3] \ 45] \ 2 \ [0 \ 1] \ 0 \ 2]$ 
990 :: now we expand opcode 2
991  $\star[\star[[[4 \ 0 \ 3] \ 45] \ 0 \ 1]$ 
992  $\star[[[4 \ 0 \ 3] \ 45] \ 0 \ 2]]$ 

```

```

993 :: mem slot 2 of the subject becomes the new formula
994 *[[[4 0 3] 45] 4 0 3]
995 **[[[4 0 3] 45] 0 3]
996 :: grab mem slot 3
997 +(45)
998 46
999

```

1000 We start above with a subject that is not a core; it's just the
 1001 atom 45. The code for `c` is then:

```
1002 [[1 4 0 3] [0 1]]
```

1003 This formula uses the Cell Maker (Distribution Rule) to insert
 1004 [4 0 3] as the head of the new subject, and puts mem
 1005 slot 1 of the old subject as the tail, so we get a new subject of
 1006 [[4 0 3] 45].

1007 What we do next is use `b` (here, 2) to select memory slot 2
 1008 from that new subject. Memory slot 2 is a formula: [4 0 3].
 1009 We run that formula against the new subject.

1010 3.5.2 Key Point/Possible Confusion

1011 When I first saw 9, I thought: “why didn’t they just use 2 to
 1012 make the transformed subject? Why is there an extra step to
 1013 use [0 1] to pull the new `*[a c]` subject? Why can’t we do
 1014 `*[a 2 c [0 b]]`?

1015 The answer is that we actually use the `*[a c]` subject *twice*:

- 1016 1. We extract from it the formula located at memory slot `b`.
- 1017 2. Then we run *that* formula against the `*[a c]` subject.

1018 If we just did `*[a 2 c [0 b]]`, then [0 b] would try to
 1019 look up the `b` memory slot in `a`, NOT `*[a c]`. So 9 gives us
 1020 one extra step to set up the core itself.

1021 3.5.3 How to Think of 9

1022 I like to think of 9 as having three parts:

- 1023 1. `c`: our formula to set up the subject, making a new sub-
 1024 ject.

- 1025 2. b: the memory slot in our new subject where an arm is.
 1026 3. Run the arm located at b against the new subject.
 1027 You can think of this as setting up a subject, pulling a stored
 1028 procedure from it, and then running that procedure against the
 1029 subject.

1030 3.6 10, Replace a Memory Slot

1031 Before explaining 10, we need to introduce a new operator, #.
 1032 # is the “edit” operator. It has the form

```
1033 #[mem-slot new-val target-tree]
```

1036 It replaces the memory slot mem-slot in target-tree with
 1037 new-val.

```
1038 :: Example
1039 #[2 [4 5] [99 88 77]]
1040 [[4 5] 88 77]
```

1041 The pseudocode for 10 is:

```
1042 *[a 10 [b c] d]    #[b *[a c] *[a d]]
```

1045 In English, this first calculates *[a c] and *[a d], and
 1046 then replaces memory slot b in the latter with the result of the
 1047 former.

1048 3.6.1 Example: Using 10 to Replace a Memory Slot

```
1049 > .*(50 [10 [2 [0 1]] [1 8 9 10]))
1050 [50 9 10]
1051
1052 :: PSEUDOCODE
1053 *[50 [10 [2 [0 1]] [1 8 9 10]]]
1054 #[2 *[50 0 1] *[50 1 8 9 10]]
1055 #[2 /[1 50] [8 9 10]]
1056 #[2 50 [8 9 10]]
1057 :: expanded as far as we can, now do the edit
1058 :: replace mem slot 2 of [8 9 10] with the value 50
1059 [50 9 10]
```

1061 We will defer discussion of opcode 11 to a later section be-
 1062 cause it is not part of strict Nock semantics.

1063 3.6.2 Towards Hoon

1064 If “grab a chunk of code from a subject and then run it against
 1065 the subject” sounds a lot like getting an arm from a core in
 1066 Hoon, that’s because it is. Nock opcode 9 is hand-in-glove
 1067 with Hoon’s core/arm structure. The “new subjects” created by
 1068 `*[a c]` are cores, and the things selected from memory slots
 1069 by `b` are arms.

1070 3.7 Real Nock Code

1071 To finish this off and take your new powers for a spin, let’s look
 1072 at some real Nock code from the wild. I got the below example
 1073 from the Dojo. It’s a mold: a function that takes a noun and
 1074 returns it if it’s the correct type, and crashes if not. The mold
 1075 here checks whether the input noun is a loobean (`0/true` or
 1076 `1/false`).

1077 3.7.1 Example: A Loobean Mold

```

1078 :: loobean mold gate
1079 > =loobean-mold ?
1080
1081
1082 :: below output truncated for our purposes here
1083 > loobean-mold
1084 < 1.fxf ... >
1085
1086 :: grab the code for the battery
1087 > -.loobean-mold
1088 [8 [6 [5 [1 0] 0 6] [1 0] 6 [5 [1 1] 0 6] [1 1] 0 0]
1089 8 [5 [0 14] 0 2] 0 6]
1090
```

1091 So we:

- 1092 1. Assign the mold gate denoted by `?` to the face `loobean-mold`.
- 1093 2. Examine what’s inside, seeing that the head is an arm.
- 1094 3. Print out that source code by calling the head.

1095 The below code acts as a formula which evaluates its sample,
1096 and returns it if it's a loobean or crashes otherwise. Let's see
1097 how that works.

1098 `[6 [5 [1 0] 0 6] [1 0] 6 [5 [1 1] 0 6] [1 1] 0 0]`

1099 We start with `6`, which means this is an if-else. The true/false
1100 test is the next element:

1101 `[5 [1 0] 0 6]`

1102 This compares the quoted value `0` (from `[1 0]`) with the value
1103 at memory slot `6` in the subject (`[0 6]`).

1104 What is the subject and what is at mem slot `6`? This code is
1105 the arm of a gate, so the subject is that gate/core: `[battery`
1106 `sample payload]`. We are looking at the battery right now,
1107 and so `[0 6]` yanks the head of the tail, or the `sample`!

1108 We know the sample will be a noun that we're testing for
1109 loobeanness, so this code starts by seeing whether the sample
1110 is the value `0`/true. If it is, the next element is `[1 0]`, so we
1111 return `0` if the sample is `0`.

1112 Otherwise, we run the second branch of the `if-else` con-
1113 ditional:

1114 `[6 [5 [1 1] 0 6] [1 1] 0 0]`

1115 This is also an opcode `6` if-else. Once again, it compares
1116 something to the value at mem slot `6` (the sample), but this
1117 time it checks whether that is the value `1`. If it is, it runs the
1118 formula `[1 1]` to return `1`.

1119 If not, it means our input was neither `0` nor `1` and is not a
1120 loobean, so we run the formula `[0 0]`, which always crashes
1121 (as there is no memory slot `0`).⁶ This is exactly what we want—
1122 crash if the sample is not a loobean!

1123 The second part of the formula which is composed with an
1124 `8` is:

1125 `[8 [5 [0 14] 0 2] 0 6]`

1126 This checks whether the battery of the current gate is equiv-
1127 alent to the head of the context; if so, it returns the sample. This

⁶Recall that in the binary tree addressing scheme of Nock, `1` refers to the entire noun, `2` to its head if it exists, and `3` to its tail if it exists.

1128 is a artifact of a type promotion in Hoon which converts a bare
1129 loobean value to a constant by pinning a constant as an exam-
1130 ple first (i.e., at slot 14). (It's completely extraneous in Nock
1131 itself, however.)

1132 3.8 Summary

1133 In this section, we've seen how to use almost all of the remain-
1134 ing opcodes, which build Nock up to a slightly more expres-
1135 sive level. We also saw how Hoon cores start to arise pretty
1136 naturally out of the Nock primitives, especially 8 and 9. Then
1137 we walked through real production Nock code to show that
1138 everything we've learned so far works exactly as expected in
1139 the wild. In the next section, you'll learn how to write real
1140 programs in Nock, and compose those programs to make new
1141 ones. Finally, hopefully you now see that, whatever its other
1142 limitations, Nock is not particularly obscurantist, and is fairly
1143 straightforward to parse, once you understand its syntax and
1144 idioms.

1145 4 Interlude

1146 In this section, we cover some loose ends deriving from the
1147 previous discussion.

1148 4.1 Order of Operations

1149 At this point, we can answer the question of what order Nock
1150 evaluates expressions in. The answer is straightforward: it
1151 starts with code that has only one pseudocode operator and
1152 all Nock code, as below:

```
1153 *[50 4 4 [5 [0 1] [1 50]]]
```

1154 The interpreter can then be thought of as moving left-to-right,
1155 expanding according to its rules, until it can't expand any fur-
1156 ther into pseudocode:

```
1157 *[50 4 4 [5 [0 1] [1 50]]]
```

```
1158 +*[50 4 [5 [0 1] [1 50]]]
```

```

1159 ++*[50 5 [0 1] [1 50]]
1160 ++=[*[50 [0 1]] *[50 [1 50]]]
1161 ++=[/[1 50] 50]

```

Once we get to that last line, ++=[/[1 50] *[50 [1 50]]], no further expansion into pseudocode is possible.

At this point, the interpreter expands “inside-out”, starting from the deepest operators. In this case, those are /[1 50] and 50, so it does those:

```

1167 ++=[50 50]

```

Then it moves to the next-furthest-inside operator: =, and then to the + operators:

```

1170 ++=[50 50]
1171 ++(0)
1172 +(1)
1173 2

```

4.1.1 Summary of Nock Order of Operations

1. On the first pass, when we have Nock code `*(nock-code)`, we expand left-to-right into pseudocode.
2. On the second pass, when we’ve fully expanded pseudocode, we evaluate operators from the inside-out.

4.2 11, Hints and Side Effects for the Interpreter

Opcode 11 lets you compute side effects and pass information to Nock’s interpreter, for that interpreter to do what it wants. The most common uses are things like jets and debugging prints. In theory, this opcode is fairly dangerous, since it tells the interpreter that it’s free to do anything.

There are two different versions of 11, depending on whether the head following 11 is an atom or a cell

```

1187 :: head is an atom (b)
1188 *[a 11 b c] *[a c]
1189
1190 :: head is a cell ([b c])
1191 *[a 11 [b c] d] *[[*[a c] *[a d]] 0 3]

```

1192 This pseudocode lets 11 handle two separate use cases:

- 1193 1. We just want the interpreter to process a message in its
1194 own language/environment.
- 1195 2. We want the interpreter to compute some Nock code
1196 with Nock semantics before processing the hint.

1197 4.2.1 Option #1: Interpreter Processes a Message

1198 Imagine the following Nock:

```
1199 . *([50 51] [11 369 0 2])
```

1200 This passes the value 369 to the interpreter, where maybe it
1201 would be the key in a hashmap. The value associated with the
1202 could be the function to debug print the whole operation. Or
1203 the value could be a jet function that specifically knows how
1204 to compute the 0 opcode faster when it's followed by a 2.

1205 After the interpreter does whatever, it needs to then discard
1206 the value 292 . 989, and computes

```
1207 > . *([50 51] [0 2])  
1208 50
```

1209 4.2.2 Option #2: Run Nock Code and then Process a Message

1210 In this case, instead of passing a static value to the interpreter,
1211 we pass it a Nock computation against the subject. The result
1212 of that computation will then be available to the interpreter to
1213 use as a message.

1214 4.2.3 Practical Use of 11: Jet Registration

1215 Jet registration in Hoon usually looks like this:

```
1216 ++ my-function  
1217 ~ / %my-function  
1218 | = var = *  
1219 ...  
1220
```

1222 ~ / “fassig” uses Nock 11 to pass the value %my-function to the
1223 interpreter, where it is associated with the C code that produces
1224 the product of my-function.

4.3 Nock in Hoon

In Hoon, you can acquire the Nock code for any function by using the `!=` “zaptis” rune. Since this produces Nock code complete with gate calls, you may also find it helpful to cast a data type as a noun using `^ - *`, which will show you the raw Nock code for that value.

As you have seen, you can also evaluate raw Nock using the `. *` “dottar” rune, which evaluates to Nock opcode 2. Other Nock runes in Hoon include `. +`, which is opcode 4; `. =`, which is opcode 5; and `. ?`, which is opcode 3. Fake opcode 12 is not actually Nock, but is used in a virtualization context to resolve values that may not be available in the current subject.

5 The Core as Design Pattern

At this point, you now know everything there is to know about Nock syntax and basic code manipulation. However, if I asked you to go write a simple program in Nock, you’d probably have a hard time getting started. It’s really helpful to see some examples of how Nock programs are designed and patterns that tend to recur in them.

We can set variables to formulas in the Urbit Dojo. We will use that in this lesson to make code easy to follow, and to show how Nock programs are made out of discrete chunks.

```
> =increment-formula [4 @ 1]
> .*(50 increment-formula)
51
> =increment-formula
```

5.1 Building a Core Manually

Every Nock program is just a formula that takes in a subject (argument). In order to run full programs against a subject, we follow the following pattern:

1. First pass: set up the dominos inside an opcode like 2 (or 7/8/9). This will create a core, and evaluate it on the second pass.

2. Second pass: evaluate that core to get our result.

The pseudocode for opcode 2, this looks like:

```
*[subject 2 core-formula call-formula]
```

Almost always, however, we'll use the 9 opcode (which is sugar on top of 2) to actually start our programs running, since it supports the concept of “cores” and “arms” very naturally. (For those who know Hoon, this “set up a core and run it” pattern should feel similar to | ^ or =>.)

5.1.1 Example: A Simple Increment Gate

Let's take code to increment the subject and turn it into a Hoon-style gate (a core with one arm and a sample). The code itself is simple: [4 0 1].

Our code expects the subject to be an atom, but Hoon gates expect their subject to be the gate itself, which has the form [battery payload]. So we need to adjust our formula to accept this subject format.

```
> =inc [7 [0 3] [4 0 1]]
:: verify that it increments value in the subject's
   tail (payload)
> .*([1 74] inc)
75
```

We put our formula in the head, and put a default (“bunted”) sample of 0 as the payload.

```
> =inc-gate [inc 0]
> inc-gate
[[7 [0 3] 4 0 1] 0]
```

Arms are best invoked using the 9 opcode. Here we use the “standard” version of it: [9 2 0 1]. Recall that 9 is $\star[a \ u \ b \ c]$, where c is the “new subject” formula, and b is the memory slot of that new subject to take an arm from. Here we use [0 1] to keep our subject (the gate) unchanged.

```
> .*(inc-gate [9 2 0 1])
1
```

```

1295     To “pass a parameter” to our gate, we use the 10 opcode to
1296     edit the subject (the inc-gate core) and replace the payload
1297     with a value. We make our subject [inc-gate val-to-increment]:

1298     :: notice how opcode 10 can replace the last element
1299       of inc-gate
1300     :: the tail of inc-gate was 0; now it's 36
1301     > .*([inc-gate 36] [10 [3 [0 3]] 0 2])
1302     [[7 [0 3] 4 0 1] 36]
1303
1304     :: now use that to set up the subject with a new
1305       tail (sample)
1306     :: and call arm in memory slot 2 (our increment
1307       formula)
1308     > .*([inc-gate 36] [9 2 10 [3 [0 3]] 0 2])
1309     37
1310     > .*([inc-gate 562] [9 2 10 [3 [0 3]] 0 2])
1311     563

```

1312 5.1.2 Summary

1313 What was the point of this? We took a perfectly good incre-
 1314 ment formula that worked on atoms, and made it more com-
 1315 plicated for the same result. *Lame!*

1316 In fact, however, this setup will come in really handy when
 1317 we have multiple arms in our core, not just one. We’ve made
 1318 it easy to set up a clean environment for increment to find its
 1319 sample in, no matter what other data is present. This will come
 1320 in *very* handy in our third example program.

1321 First, however, let’s look at putting a more complicated for-
 1322 mula inside a gate. (Hint: the process is the same).

1323 5.1.3 Example: A Decrement Gate

1324 Here we take a piece of code that decrements the subject when
 1325 the subject is an atom, and turn it into a gate. For now, we
 1326 will take it as a given that the decrement function given below
 1327 works. (The code will crash when the input is a cell or `0`)

```

1328 :: the code: set a Dojo face =dec
1329

```



```

1330 > =dec [6 [5 [0 1] [1 0]] [0 0] [6 [3 0 1] [0 0] [8
1331      [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0
1332      6] 0 7] 9 2 0 1]]]
1333
1334 :: testing in the Dojo
1335 > .* (100 dec)
1336 99
1337
1338 :: with non-atom input we crash
1339 .* ([100 101] dec)
1340 dojo: hoon expression failed
1341

```

Let's re-jigger the decrement formula to accept the subject in format [battery payload]:

```

1344 :: gets value in mem slot 3 and applies the
1345      decrement formula to it
1346 > =dec-arm [7 [0 3] dec]
1347
1348 :: verify that dec-arm decrements the value in the
1349      tail
1350 > .* ([1 88] dec-arm)
1351 87
1352
1353 :: output the full dec-arm formula code
1354 > dec-arm
1355 [7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1356      0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1357      0 7] 9 2 0 1]

```

Next, assemble the core:

```

1359 :: all we need to do is add the sample! (we give it a
1360      default value of 0)
1361 > =dec-gate [dec-arm 0]
1362 > dec-gate
1363 [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1364      0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1365      0 7] 9 2 0 1] 0]
1366 [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1367      0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1368      0 7] 9 2 0 1] 0]

```

```

1369     As in the first example, we use the 10 opcode to edit the
1370     subject (the dec-gate core) and replace the payload with a
1371     value. We make our subject [dec-gate val-to-decrement].

1372     :: confirm that our 10 code places 36 in the tail
1373     > .*([dec-gate 36] [10 [3 [0 3]] 0 2])
1374     [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1375       0 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1376       0 7] 9 2 0 1] 36]
1377
1378     :: now use that to set up the subject with a new
1379     tail (sample)
1380     :: and call arm in memory slot 2 (our decrement
1381     formula)
1382     > .*([dec-gate 36] [9 2 10 [3 [0 3]] 0 2])
1383     35
1384     > .*([dec-gate 562] [9 2 10 [3 [0 3]] 0 2])
1385     561
    
```

1386 5.1.4 Example: Comparing Two Numbers

1387 In this example we will build a Nock core that compares two
 1388 numbers and returns whether they are greater than, less than,
 1389 or equal to each other.

1390 The program specification is to take two numbers `a` and `b`
 1391 and return:

- 1392 • 0 if `a == b`.
- 1393 • 1 if `a > b`.
- 1394 • 2 if `a < b`.

1395 The trick is that the only mathematical operators we have
 1396 for this are the 4 opcode for increment, 5 for equals, and a
 1397 decrement Nock function that we will supply. The basic algo-
 1398 rithm is:

- 1399 • If `a == b`, return 0.
- 1400 • If `a == 0`, return 2. (I.e, `a < b`).
- 1401 • If `b == 0`, return 1. (I.e., `a > b`).

- Recur with a and b both decremented.

The above works because if a is smaller than b, it hits 0 first. If b is smaller than a, it hits 0 first. If they are equal, we never decrement in the first place, so there are no issues with numbers going negative.

We will construct a core that has two arms: one arm with the algorithm logic, and a second arm with the decrement gate from the second example. These arms will live in the head of the core (think “battery”).

We will put our variables a and b in the tail of the core (think “payload”). They will be updated before the core is called each time. (This is similar to a trap or door in Hoon).

The final core we will create will have the structure [battery payload], which can be broken down as:

```
[[main-logic dec-gate] a b]
```

Once the core is set up, we’ll invoke the algorithm logic arm to set it running. We will build the main logic “inside-out”, by first defining the “recur” code, and then inserting it into our if/else tests that implement the algorithm.

This recursion logic assumes the decrement gate lives in the tail of the battery (the battery is the head), i.e. [0 5]. We are using [10 [3 [0 memory-slot]] 0 5] to extract the decrement gate from our core, and then we invoke it with [9 2 ...].

This is identical to what we did in the second example—the only difference is that we now do it twice: once for memory slot 6 (a) and again for memory slot 7 (b).

```
> =dec-a [9 2 10 [3 [0 6]] 0 5]
> =dec-b [9 2 10 [3 [0 7]] 0 5]

:: we can test our formulas by making a
:: dummy core (a = 33, b = 77)
:: first formula in battery just returns
:: the current subject
> =dummy-core [[[0 1] dec-gate] 33 77]
> .* (dummy-core dec-a)
32
```

```

1439 > .*(dummy-core dec-b)
1440 76

```

1441 The recursion formula here operates by invoking its subject—
 1442 itself—and applying the `dec-a` and `dec-b` formulas to it. Then
 1443 we use opcode 9 for the new subject's setup: replace payload
 1444 with new `a` and `b` run the current core

```

1445 > =recur [9 4 [[0 2] dec-a dec-b]]
1446
1447 :: test it -- we want to see same battery,
1448 :: with payload being replaced with 32 and 76
1449 > .*(dummy-core recur)
1450 [[[0 1] [7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0
1451      0] 8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2]
1452      [4 0 6] 0 7] 9 2 0 1] 0] 32 76]
1453
1454 :: clean up the dummy-core
1455 > =dummy-core

```

1456 The main logic is then:

```

1457
1458 :: check whether a is 0; return 2 if true
1459 :: else check whether b is 0; return 1 if true
1460 :: else recur
1461 > =main-logic [6 [5 [0 6] [1 0]] [1 2] [6 [5 [0 7] [1
1462      0]] [1 1] recur]]
1463
1464 :: test (we aren't doing the equality case here)
1465 :: a < b
1466 > .*([[main-logic dec-gate] 9 10] [9 4 0 1])
1467 2
1468 :: a > b
1469 > .*([[main-logic dec-gate] 10 9] [9 4 0 1])
1470 1
1471

```

1472 We build an outer test for whether `a==b`, and then if it's
 1473 not, we use the 1 opcode to return the core.

```

1474 > =battery [main-logic dec-gate]
1475
1476 :: payload for our core is memory slots 2 and 3 (a
1477 and b)
1478 > =comparison [6 [5 [0 2] [0 3]] [1 0] [9 4 [[1
1479      battery] [0 2] 0 3]]]

```

```

1480
1481 :: test our 3 cases
1482 > .*([0 8] comparison)
1483 2
1484 > .*([0 0] comparison)
1485 0
1486 > .*([8 0] comparison)
1487 1

```

1488 We can take this code and use it *anywhere* that our subject
 1489 is a cell of two numbers, and it will “just work”.

```

1490 [6 [5 [0 2] 0 3] [1 0] 9 4 [1 [6 [5 [0 6] 1 0] [1 2]
1491     6 [5 [0 7] 1 0] [1 1] 9 4 [0 2] [9 2 10 [3 0 6] 0
1492     5] 9 2 10 [3 0 7] 0 5] [7 [0 3] 6 [5 [0 1] 1 0]
1493     [0 0] 6 [3 0 1] [0 0] 8 [1 0] 8 [1 6 [5 [0 7] 4 0
1494     6] [0 6] 9 2 [0 2] [4 0 6] 0 7] 9 2 0 1] 0] [0 2]
1495     0 3]

```

1496 One common convention used in languages like Hoon is
 1497 to denote Nock rules as constants, e.g. %0 (corresponding to
 1498 our 0). This makes it much more straightforward to interpret
 1499 hand-rolled Nock code with addresses and constants present.

1500 6 Conclusion

1501 After this tutorial, you should be prepared to interpret Nock
 1502 code when it is produced by Hoon. Although there is only
 1503 rarely any call to write Nock code directly, you should now
 1504 have a good understanding of how to read and interpret it. As
 1505 you write your own interpreters in the future, you may even
 1506 have call to hand-roll Nock expressions on occasion.