

1

2

3

4

5
6
7
8
9
10
11
12
13
14
15
16
17
18
19

20

21
22
23
24
25
26

Nock for Everyday Coders

Tim Galebach ~timluc-miptev
Hyperware

Abstract

Nock is not super complex, and most programmers can learn the basics of it rapidly. The mental model gained by learning Nock turns out to be very useful in learning Hoon and understanding Urbit. While the Urbit docs generally suggest to not worry about Nock, Nock is very simple and small. Most programmers will feel more comfortable in Hoon having learned Nock’s basics. The goal of this tutorial is to explain Nock clearly in terms most programmers will relate to, to impart a feeling of confidence with very basic Nock, and to give a knowledge of Nock’s idioms and big wins so that they carry over to learning Hoon.

This tutorial was originally published online by ~timluc-miptev in early 2020 and remains an excellent exposition of Nock’s affordances.

Contents

1	Getting Started	2
1.1	Phases of Nock	3
1.2	What Is a Nock Interpreter?	3
1.3	How to Run Nock Code	4
1.4	Subject, Formula?	4
1.5	Evaluating Our First Nock Code	4

27	2 Fundamental Opcodes	5
28	2.1 Nock's Simplest Functions, 0 and 1	5
29	2.2 Binary Tree Addressing	5
30	2.3 0, the "Memory Slot" Function	7
31	2.4 1, the "Quoter" Function	9
32	2.5 4, the Incrementing Function	10
33	2.6 The Cell-Maker (AKA the Distribution Rule) . .	13
34	2.7 3, the Cell Detector, and 5, the Equality Tester	15
35	2.8 2, the "Subject-Altering" Function	19
36	2.9 Summary of Fundamental Opcodes	22
37	3 Composite Opcodes	23
38	3.1 An Aside About the 1 (Quoter Function) . . .	23
39	3.2 6, "If/Else" Conditional Branch	24
40	3.3 7, the "Composition" Opcode	28
41	3.4 8, the "Variable Push" Opcode	28
42	3.5 9, Run a Stored Procedure Arm in a Core . . .	30
43	3.6 10, Replace a Memory Slot	32
44	3.7 Real Nock Code	33
45	3.8 Summary	35
46	4 Interlude	35
47	4.1 Order of Operations	35
48	4.2 11, Hints and Side Effects for the Interpreter .	36
49	4.3 Nock in Hoon	38
50	5 The Core as Design Pattern	38
51	5.1 Building a Core Manually	38
52	6 Conclusion	45

53 1 Getting Started

54 When people first look at Nock, they see the definition, which
 55 is fairly intimidating. I'm talking about lines like this:

56 `/[(a + a) b]` `/[2 /[a b]]`
 57

59 The problem is, the programmer may already know that Nock
60 code looks more like the below – just lists of numbers, with no
61 symbols:

```
62 [6
63  [5 [0 6] [1 0]]
64  [0 7]
65  [9 4 [[0 2] [2 [0 6] [0 5]] [4 0 7]]]
66 ]
```

67 What gives? Which one is the “real” Nock?

68 1.1 Phases of Nock

69 We are looking at two different things in the examples above:

- 70 1. Pseudocode for *how to interpret* Nock.
- 71 2. Code to be interpreted (written as lists of numbers).

72 The symbols and spec are pseudocode, not real Nock code.
73 They could just as easily be written in English, and they will
74 never be written down as actual Nock code and given to an
75 interpreter. They represent what an interpreter should do to
76 turn Nock code into interpreter instructions.

77 The lists of numbers are the actual Nock code. This is what
78 you feed to an interpreter to get some result.

79 1.2 What Is a Nock Interpreter?

80 An interpreter can be a computer program, or it can be a human
81 manually expanding Nock code into results. In both cases, the
82 program and human have to know the Nock pseudocode in
83 order to do the right thing with incoming Nock code.

84 So a Nock interpreter is any entity that takes Nock code as
85 input, and gives a noun as output. A noun can be:

```
86 :: a number
87 782
88 :: a cell (pair with two elements)
89 [782 9872]
90 :: each element can itself be a pair
```

```

91 [782 [9872 89728]]
92 :: the above can be written as
93 [782 9872 89728]
```

94 1.3 How to Run Nock Code

95 We will be expanding Nock pseudocode manually in the exam-
 96 ples that follow, in effect acting as our own interpreter.

97 If we want to check that we’re getting the right results from
 98 our manual interpretation, we need to run a Nock interpreter,
 99 such as the Urbit Dojo.

- 100 1. Start up a Dojo session on a fake ship.¹
- 101 2. At the prompt, we can execute Nock using `. * dottar`, e.g.,
 102 `. * (NOCK_SUBJECT, NOCK_FORMULA).`

103 1.4 Subject, Formula?

104 Let’s keep this simple:

- 105 • subject = an argument to a function
- 106 • formula = the function

107 That’s it. We’ll see below how this works, going really
 108 slowly with examples.

109 1.5 Evaluating Our First Nock Code

110 OK, so the interpreter takes two arguments, a “subject” and a
 111 “formula”. Both are nouns (a number or a cell). Let’s run some
 112 insanely simple Nock code in the Dojo:²

```

113 > . * (42 [0 1])
114 42
```

¹Consult <https://docs.urbit.org> for details.

²Code samples beginning with `>` are Dojo inputs (to wit, Hoon expressions).

115 In the above, 42 is our subject. [0 1] is our formula.
 116 Formulas are always cells, and the first element of the cell
 117 is a number that you can think of as *the name of the function*.
 118 In this case, our function name is 0, which is the memory
 119 slot function. It is always followed by 1 number, in this case 1,
 120 which is the number of the memory slot to fetch in the subject.
 121 Whenever we look at Nock code, we want to ask:

- 122 • What is the subject (function argument)? In this case,
 123 it's 42.
- 124 • What is the formula (function)? In this case, it's [0 1].
- 125 • What value does that formula (function) produce when
 126 called on this subject (argument)? In this case, the return
 127 value is 42.

128 Why is the return value 42? How does this formula work?

129 2 Fundamental Opcodes

130 2.1 Nock's Simplest Functions, 0 and 1

131 The two most basic Nock functions are 0 address and 1 con-
 132 stant. The goal here is to get strong intuitions of what they do,
 133 how they handle edge cases, and how this relates to the Nock
 134 spec/pseudocode.

135 2.2 Binary Tree Addressing

136 Before getting started on Nock proper, we should understand
 137 how Nock and Hoon handle addresses in binary trees. If you
 138 already understand why memory slot 5 of [['apple' %pie]
 139 [0b1101 0xdad]] is %pie, you are good to go and can skip
 140 ahead to the Section 2.3.

141 Every noun in Nock can be thought of as a tree, which
 142 means we can give an exact number to access any position in
 143 the tree. This means that, no matter how big our subject (ar-
 144 gument) is, we can yank a value out of any part of it.

145 The Nock specification defines noun tree addressing:


```

146
147 / [1 a]           a
148 / [2 a b]        a
149 / [3 a b]        b
150 / [(a + a) b]     / [2 / [a b]]
151 / [(a + a + 1) b] / [3 / [a b]]
152

```

153 This permits the address to be defined within a particular sub-
154 tree as well as within the overall tree (Figure 1).

Figure 1: A binary tree with labeled node addresses to several layers.



155 That is, how do you say which slot number you want from
156 a given tree? We say that:

- 157 • The tree root is address 1.
- 158 • The head of every node n is $2n$.
- 159 • The tail of every node n is $2n + 1$.


160 Let's take an example tree to illustrate. In Nock cell form,
161 the tree is:

```
162 [[4 5] [6 14 15]]
```

163 Diagrammatically, the tree looks like Figure 2.

- 164 • 1 is the address of the whole tree, `[[4 5] [6 14 15]]`.
- 165 • 2 is the address of the left branch, `[4 5]`.

Figure 2: A binary tree with some labeled node addresses.



- 3 is the address of the right branch, [6 14 15].
- 15 is the value 15.

Let's play around now in the Dojo Nock interpreter so that we can confirm this. In each example, our subject (argument) will be the tree [[4 5] [6 14 15]].

```

171 :: formula (function) [0 1]: get the whole tree
172 > .*([[4 5] [6 14 15]] [0 1])
173 [[4 5] 6 14 15]
174
175
176 :: formula (function) [0 2]: get the left branch
177 > .*([[4 5] [6 14 15]] [0 2])
178 [4 5]
179
180 :: formula (function) [0 7]: get the subtree in slot 7
181 > .*([[4 5] [6 14 15]] [0 7])
182 [14 15]
183

```

2.3 0, the "Memory Slot" Function

The pseudocode for the 0 opcode³ is as follows:

```

186 * [a 0 b] / [b a]
187
188

```

³From this point on, we will write Nock opcodes in special type so that they can be distinguished from other numbers in Nock. This reflects Hoon practice, which marks Nock rules as constants like %0. While not strictly necessary, it can be helpful when learning or coding Nock to see an explicit distinction.

189 /[b a] is pseudocode. In English, it means “a is a binary
190 tree. Get the memory slot numbered b.

191 So if a were the tree [9 10], and b were 1, we’d get the
192 memory slot 1 in the tree [9 10], which is just the tree itself.

193 Written in pseudocode, that’s /[1 [9 10]]. We can also
194 do /[2 [9 10]], which grabs memory slot 2, aka 9.

195 Let’s look at some examples. I’ve put them first in Dojo
196 form so that you can see how they run in that interpreter, and
197 then I show the “human interpreter” in pseudocode below that.

198 2.3.1 Example: Retrieve the Subject

```
199 :: get memory slot 1
200 > .*([50 51] [_ 1])
201 [50 51]
202
203 :: PSEUDOCODE -- replaces the Dojo's () with []
204 *[[50 51] [_ 1]]
205 :: *[a @0@ b] -> a = [50 51], b = 1
206 /[1 [50 51]]
207 [50 51]
```

209 2.3.2 Example: Retrieve the Head of the Subject

```
210 > .*([50 51] [_ 2])
211 50
212
213 :: PSEUDOCODE -- replaces the Dojo's () with []
214 *[[50 51] [_ 2]]
215 :: *[a @0@ b] -> a = [50 51], b = 2
216 /[2 [50 51]]
217 50
```

219 2.3.3 Example: A Crash

```
220 > .*([50 51] [_ [0 1]])
221 dojo: hoon expression failed
222 :: PSEUDOCODE
223 *[[50 51] [_ [0 1]]]
224 /[[0 1] [50 51]]
```



```

226 :: can't evaluate this--a memory slot must be a
227     number like 2, not a cell like [0 1]
228 CRASH
229

```

230 2.3.4 Summary of @

231 @ is everywhere in Nock, because “get something from a mem-
 232 ory slot” really means “store stuff in a place and get it whenever
 233 I want,” which is another name for creating variables. These
 234 memory slot numbers take the place of variable names.

235 If you’re familiar with assembly or C, they’re conceptu-
 236 ally similar to memory pointers or registers in how Nock uses
 237 them. Keep in mind though, Nock is functional/immutable, so
 238 it doesn’t update memory locations: it creates copies of data
 239 structures with altered values.

240 We’ve also seen that the memory slot function can’t take
 241 just anything as the memory slot to fetch: it must receive an
 242 atom (number).

243 2.4 1, the “Quoter” Function

244 This is another really simple function. You can think of it as a
 245 “quoter”: it just returns any value passed to it exactly as it is.
 246 It ignores the subject, and just quotes the value after it. Let’s
 247 look at a couple examples.

```

248 > .*([20 30] [1 67])
249 67
250 :: [@1@ 2 587] is same as [@1@ [2 [587]]]
251 > .*([20 30] [1 [2 [587]]]
252 [2 587])

```

253 It doesn’t matter how much information is after the 1: 1 is
 254 a dumb function that just returns it all.

255 The pseudocode for 1 is:

```

256 * [a 1 b]    b
257
258

```

259 In English, this means: “ignore the subject ‘a’, and just re-
 260 turn everything after the 1 exactly as it is.

261 Let's look at our first example, `.*([20 30] [1 67])`. The
 262 subject `a` is `[20 30]`, so we ignore that. What's after the `1`?
 263 `67`, so we return that.

264 In the second example, "everything after the `1`" is longer,
 265 but the same rule applies: the Nock interpreter just returns it
 266 exactly as it is, after stripping out unnecessary brackets.

267 2.4.1 Summary of `1`

268 `1` is a simple function that just returns whatever is after it
 269 ("quoter" or "constant"). It's useful for quoting values that you
 270 want to use later in your Nock code.

271 2.5 `4`, the Incrementing Function

272 `0` and `1` are simple functions that don't have any nested be-
 273 havior. Now we're going to move to a function that *does* have
 274 nested children, opcode `4`. In these examples, pay attention to
 275 how `4` operates on its arguments; we'll look at pseudocode and
 276 break down the examples in a moment.

```
277 > .*(50 [4 0 1])
278 51
279
280 > .*(50 [4 4 0 1])
281 52
282
283 > .*([100 150] [4 4 0 3])
284 152
285
286 > .*(50 [4 1 98])
287 99
288
289 > .*(50 [4 1 [0 2]])
290 dojo: hoon expression failed
```

291 Here's `4`'s pseudocode, juxtaposed with that of `0` and `1`:

```
292 *[a 4 b]   +*[a b]
293
294 *[a 0 b]   /[b a]
295
296 *[a 1 b]   b
```

297 In English, this says “when we have subject a, function 4,
 298 and Nock code b, first evaluate [a b] as [subject formula],
 299 and then add 1 to the result.”

300 If we contrast with 0 and 1, we see that the right side has
 301 a \star symbol. This symbol means “evaluate the expression again
 302 in the Nock interpreter.” 0 and 1 *did not* have this symbol, and
 303 that’s why they couldn’t evaluate nested Nock expressions.

304 2.5.1 Example: Walking Through Increment

305 Let’s start by translating the Dojo’s $\star(\text{subject formula})$ to
 306 pseudocode of the form $\star[a\ b]$, and then expand it line by
 307 line:

```
308 > .*(50 [4 0 1])
309 51
310 :: change to pseudocode
311 *[50 [4 0 1]]
312 :: move the 4 to the outside as +
313 **[50 [0 1]]
314 :: expand the 0 opcode to the memory slot operator "/"
315 +/[1 50]
316 :: grabs the memory slot
317 +(50)
318 :: evaluate
319 51
```

320 The [a b] part of $\star[a\ b]$ expands to:

321 [50 [0 1]]

322 OK, this we know how to handle! It’s just our 0 function, and
 323 it wants the value in memory slot 1 of the subject. That’s 50.

324 Now we know that $\star[a\ b]=50$, and we just have to add 1
 325 to it (the + in $\star[a\ b]$). That is 51, which is exactly what the
 326 interpreter gave us.

327 2.5.2 Example: Walking Through Increment

328 This one is similar, we just have an extra 4. We again start by
 329 translating the Dojo’s $\star(\text{subject formula})$ to pseudocode,
 330 and then expand

```

331 > .* (50 [4 4 0 1])
332 52
333 *[50 [4 4 0 1]]
334 :: the first 4 moves outside as a '+'
335 **[50 [4 0 1]]
336 :: 2nd 4 becomes a '+'...we have opcode 0 again!
337 ***[50 [0 1]]
338 ++/[1 50]
339 ++(50)
340 +(51)
341 52

```

342 2.5.3 Example: Walking Through Serial Increment

343 Here we again see lots of 4s applied consecutively, and we also
 344 see how we can yank values out of a more complicated subject
 345 and manipulate them. Notice how the subject is a cell, not an
 346 atom. The rest is the same as in the previous example.

```

347 > .* ([100 150] [4 4 0 3])
348 152
349 a (the subject) = [100 150]
350 *[[100 150] [4 4 0 3]]
351 **[[100 150] [4 0 3]]
352 :: Now we've extracted all the increments,
353 :: so we just grab the value at memory slot 3
354 ***[[100 150] [0 3]]
355 ++/[3 [100 150]]
356 ++(150)
357 +(151)
358 152

```

359 2.5.4 Example: Walking Through Incrementing a Constant

360 In the below example, we see how we can use the quote/con-
 361 stant function 1 to generate the value 98 and increment it. We
 362 ignore the subject 50 completely.

```

363 > .* (50 [4 1 98])
364 99
365 *[50 [4 1 98]]
366 :: formula is the "quoter" function

```

```

367 ++[50 [1 98]]
368 +(98)
369 99

```

370 2.5.5 Example: A Crash When Incrementing a Cell

371 Just as opcode 0 had values it couldn't handle (non-atoms), so
 372 opcode 4 needs the nested value inside it to evaluate to an atom.

```

373 > .* (50 [4 1 [0 2]])
374 dojo: hoon expression failed
375 *[50 [4 1 [0 2]]]
376 :: OK cool, the nested value is a 1 opcode
377 ++[50 [1 [0 2]]]
378 :: 1 ignores the subject (50) and just returns [0 2]
379 +([0 2])
380 :: [0 2] isn't an atom, so how can we increment it???
381 :: We can't, so we crash.
382 dojo: hoon expression failed

```

383 2.5.6 Summary of 4

384 In these examples, we've seen that function 4 can be called as
 385 many times in a row as we want. At the end of those calls, it
 386 always ends up incrementing a number that either is yanked
 387 from the subject (memory slot function 0) or quoted as it is
 388 (quote function 1).

389 2.6 The Cell-Maker (aka the Distribution Rule)

390 The Nock interpreter is allowed to return nouns, which are
 391 atoms (positive numbers) or cells (pairs of nouns). What have
 392 our functions/opcodes been returning so far?

- 393 • 0: atoms or cells, depending what's in the memory slot
 394 that we yoink.
- 395 • 1: atoms or cells, depending on what we quote.
- 396 • 4: just atoms.

But what if my subject was [51 67 89], and I wanted to increment every value and return that as [52 68 90]? How can I do that when it's a cell, and 4 only seems able to return atoms?

The answer is something that the Nock docs call the “distribution rule” or “implicit cons” (hello, fellow LISPer!), but that I find easiest to think of as the “Cell-Maker Rule”.

2.6.1 A Quick Detour into Nock Formulas

We haven't talked much yet about what values are legal to feed into the Nock interpreter (the `.*(subject formula)` function in the Dojo). So far, we've only been using formulas that start with numbers (our functions/opcodes 0/1/4).

Let's now fully solidify our understanding of what's a legal formula by taking a quick look at the three possible cases. (The format is `*(subject formula)`.)

- The formula is a cell starting with an opcode. We know this is OK.

```
> .*(50 [0 1])
50
```

- The formula is just an atom (o). This is not valid.

```
> .*(50 0)
dojo: hoon expression failed
```

- Finally, we try a formula that is a cell starting with an atom. This looks invalid, but—it works!

```
> .*(50 [[0 1] [1 203]])
[50 203]
```

So apparently a formula cell can start with a cell. The Cell-Maker rule is:

```
*[subject [formula-x formula-y]]→
  [*[subject formula-x] *[subject formula-y]]
```

429 In our example above, `formula-x` is `[0 1]`, and `formula-y`
 430 is `[1 203]`. They each evaluate individually against the sub-
 431 ject, and the end result is a cell.

432 We can make as many cells in a row as we want:

```
433 > .* (50 [[0 1] [1 203] [0 1] [1 19] [1 76]])
434 [50 203 50 19 76]
```

435 We can put any operation inside each cell:

```
436 > .* ([19 20] [[0 1] [1 76] [4 4 0 3]])
437 [[19 20] 76 22]
```

438 If we take the returned collection `[[19 20] 76 22]` in or-
 439 der, we can write in English how they connect to our collection
 440 of formulas that we passed:

- 441 • `[0 1]`: get memory slot 1
- 442 • `[1 76]`: return the quoted value 76
- 443 • `[4 4 0 3]`: increment twice the value in memory slot 3
 444 (20)

445 So we can pass one small subject `([19 20])` and make an
 446 arbitrarily long collection of values from it, using any functions
 447 we want. Cell-Maker FTW!

448 2.7 3, the Cell Detector, and 5, the Equality Tester

449 Now we come to functions/opcodes 3 and 5, which are pretty
 450 straightforward after we've seen how 4 and the Cell-Maker
 451 work. Functions 3 and 5, like 4, allow nested evaluation. Let's
 452 put all their pseudocode definitions together to compare:

```
453 :: function/opcode 3
454 * [a 3 b]    ?*[a b]
455 :: function/opcode 4
456 * [a 4 b]    +*[a b]
457 :: function/opcode 5
458 * [a 5 b c]   =*[a b] *[a c]
```

461 First of all, notice how the right side of all these “equations”
 462 has the evaluation operator, `*`. This means that these functions
 463 can have nested formulas, since they keep evaluating all the
 464 way down.

465 There are some new pseudocode symbols here that we need
 466 to translate into English. We already know `+`: “increment the
 467 value after this”. Now we also see:

- 468 • `?:` “check whether the value after this is a cell. Return 0
 469 if yes, 1 if no”.⁴
- 470 • `=:` “first run the function in b with subject a as the argu-
 471 ment, and same for the function in c. If the results are
 472 equal, return 0, if not, return 1.

473 2.7.1 Example: Not a Cell

```
474 > .*(50 [3 0 1])
475 1
476
477 :: PSEUDOCODE OF THE ABOVE
478 *[50 [3 0 1]]
479 ?*[50 [0 1]]
480 ::get memory slot 1 of the subject
481 ?(50)
482 :: is 50 a cell? No, so return 1
483 1
484
```

485 2.7.2 Example: A Cell

```
486 > .*([50 51] [3 0 1])
487 0
488
489 :: PSEUDOCODE OF THE ABOVE
490 *[[50 51] [3 0 1]]
491 ?*[[50 51] [0 1]]
492 ::get memory slot 1 of the subject: [50 51]
493 ?([50 51])
494 :: is [50 51] a cell? Yes, so return 0
495 0
496
```

⁴Nock conventionally explains that there is one “true” case and (potentially) many “false” cases, thus `0/true` and `1/false`.

2.7.3 Example: Nested Cell Evaluation with 4

```

497
498 > .*([50 51] [4 4 3 0 1])
499 2
500
501 :: PSEUDOCODE OF THE ABOVE
502 *[[50 51] [4 4 3 0 1]]
503 **[[50 51] [4 3 0 1]]
504 :: whatever comes out of the 3 function,
505 :: we're gonna increment it twice
506 ***[[50 51] [3 0 1]]
507 :: down to just fetching memory slot 1
508 ++*[[50 51] [0 1]]
509 ++?([50 51])
510 :: is [50 51] a cell? Yes, so return 0
511 ++(0)
512 +(1)
513 2
514

```

2.7.4 Example: Check Multiple Cases of Cells

```

516
517 > .*([50 51] 52) [[3 0 2] [3 0 3]]
518 [0 1]
519 *[[50 51] 52] [[3 0 2] [3 0 3]]
520 ?*[50 51] 52] [0 2]]
521    *[50 51] [0 3]]
522 :: yank memory slots 2 and 3
523 ?*[50 51] 52]
524 :: first is a cell, second is not
525 [0 1]
526

```

2.7.5 Example: Equality Test

Because 5 compares the results of 2 formulas, it *always* makes 2 inner evaluations of the subject. It's similar to Cell-Maker in this way.

```

531
532 > .*([50 51] [5 [0 2] [0 2]])
533 0
534 :: PSEUDOCODE
535 *[[50 51] [5 [0 2] [0 2]]]

```

```
536 :: factor out the =
537 =[*[[50 51] [0 2]] *[[50 51] [0 2]]]
538 :: get memory slot 2 twice
539 =(50 50)
540 0
541
```

542 2.76 Example: Comparing Two Unequal Values

```
543
544 > .*([50 51] [5 [0 2] [0 3]])
545 1
546 :: PSEUDOCODE
547 *[[50 51] [5 [0 2] [0 3]]]
548 :: factor out the =
549 =[*[[50 51] [0 2]] *[[50 51] [0 3]]]
550 :: get memory slot 2 and memory slot 3
551 =(50 51)
552 1
553
```

554 2.77 Example: Compare Values with Function Calls

```
555
556 > .*([50 51] [5 [4 0 2] [0 3]])
557 0
558 :: PSEUDOCODE
559 *[[50 51] [5 [4 0 2] [0 3]]]
560 :: factor out the =
561 =[*[[50 51] [4 0 2]]
562   *[[50 51] [0 3]]]
563 :: factor out the +
564 =[+*[[50 51] [0 2]]
565   *[[50 51] [0 3]]]
566 :: get memory slots 2 and 3
567 =[+50 51]
568 :: evaluate the +
569 =(51 51)
570 0
571
```

572 2.78 Example: Compare Cells with Function Calls

```

573
574 > . * ([99 99] [5 [1 [99 99]] [0 1]])
575 0
576 :: PSEUDOCODE
577 * [[99 99] [5 [1 [99 99]] [0 1]]]
578 = [* [[99 99] [1 [99 99]]]
579    * [[99 99] [0 1]]]
580 :: first cell is the result of quoter function
581 :: second cell is the result of memory fetch
582 = [[99 99] [99 99]]
583 :: true
584 0
585

```

586 2.7.9 Summary of 3 and 5

587 These along with 4 are known as the “axiomatic” functions.
588 They exist to implement definitional logic for the Nock spec-
589 ification. In particular, the cell check and the equality check
590 provide for structural analysis of nouns.

591 2.8 2, the “Subject-Altering” Function

592 In all our examples so far, the subject has been defined at the
593 start when we call the interpreter, and never changes. But what
594 if we want to change the subject?

595 A different subject? Why would we want that? Here’s an
596 easy example. Say you found the following piece of Nock code
597 on the interwebz:

```

598 [8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2]
599    [4 0 6] 0 7] 9 2 0 1]

```

600 This is the code for a Nock function that expects a subject
601 that is an atom, and decrements that subject by 1. You can
602 actually enter it in the Dojo right now:

```

603 > . * (100 [8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0
604    2] [4 0 6] 0 7] 9 2 0 1])
605 99

```

606 This works! You do not have to understand the code right now;
607 just try entering different numbers instead of 100 to see the
608 program working.

609 But now imagine that I have a Nock program with a differ-
 610 ent subject

611 > .*([50 51] some-formula)

612 And somewhere inside `some-formula`, I want to decre-
 613 ment the number 51. I can't pass `[50 51]` as the subject to
 614 my decrementer code above though, because that's a cell, and
 615 it expects just an atom (a number).

616 2.8.1 Subject Altering to the Rescue

617 The function/opcode `2` is designed to handle this problem for
 618 us. First the pseudocode:

619
 620 `*(a 2 b c) *[*[a b] *[a c]]`
 621

622 `2` expects 2 formulas after the subject `a`: `b` and `c`. With
 623 those, it:

- 624 1. runs formula `b` against the subject to set up a new envi-
 625 ronment/subject derived from the subject
- 626 2. runs formula `c` against the subject to prepare a 2nd func-
 627 tion
- 628 3. run that 2nd function against the new environment/sub-
 629 ject from step (1)

630 Note that the pseudocode for `2` has nested “`*`”s.

631 `*([*[a b] *[a c]]`

632 The two inner `*s` run steps (1) and (2), and the outer one,
 633 around the whole expression, runs step (3).

634 2.8.2 Example: Change the Subject and Call a Constant Value

635
 636 > .*([50 51] [2 [0 3] [1 [4 0 1]]])
 637 52
 638 :: PSEUDOCODE
 639 `*([50 51] [2 [0 3] [1 [4 0 1]]])`
 640 :: separate `b` and `c` to each run against the subject
 641 (steps 1 and 2)

```

642 *[*[[50 51] [0 3]] *[[50 51] [1 [4 0 1]]]]
643 :: after steps 1 and 2, we have a new subject, 51!
644 :: note how we're back in normal *[subject formula]
645     form
646 *[[51 [4 0 1]]
647 :: apply the 4 function as we're used to
648 **[[51 [0 1]]
649 :: grab 51 from memory slot 1
650 +(51)
651 52
652

```

653 2.8.3 Example: Grab a Block of Code from the Subject and Run It

Think of this as grabbing a “stored procedure” from the subject.

```

654
655 > .*([[4 0 1] 51] [2 [0 3] [0 2]]])
656 52
657
658 :: PSEUDOCODE, subject is [[4 0 1] 51]
659 *[[[4 0 1] 51] [2 [0 3] [0 2]]]
660 *[*[[[4 0 1] 51] [0 3]
661     *[[[4 0 1] 51] [0 2]]]]
662 :: step 1 gets memory slot 3, step 2 grabs memory
663     slot 2
664 *[[51 [4 0 1]]
665 :: looks like a normal 4 opcode to me!
666 **[[51 [0 1]]
667 :: grab memory slot 1
668 +(51)
669 52
670

```

671 2.8.4 Example: Back to Our Motivating Case

Remember our decrementing block of code that we couldn’t use when the subject was [50 51], instead of just an atom? Opcode 2 makes handling that issue a piece of cake.

We simply use 2 to transform our subject into an atom, and use 1 to quote the decrement block of code before it evaluates in step (3).

```

678
679 > .*([50 51] [2 [0 2] [1 [8 [1 0] 8 [1 6 [5 [0 7] 4 0
680     6] [0 6] 9 2 [0 2] [4 0 6] 0 7] 9 2 0 1]]])

```

```

681 49
682 :: PSEUDOCODE (subject is [50 51])
683 :: ("decrement-formula" substituted for clarity)
684 *[[50 51] [2 [0 2] [1 decrement-formula]]]
685 *[*[[50 51] [0 2]] *[[50 51] [1 decrement-formula]]]
686 *[[50 decrement-formula]
687 ::decrement-formula has the atom subject that it wants
688 49
689

```

690 2.8.5 Summary of 2

691 For those who know a bit of Hoon, the second example above
692 is rather similar to calling an arm that produces a gate, and
693 then running the gate. Most of Hoon runs this type of stored-
694 procedure plus subject-altering Nock, and it all uses opcode 2
695 at its base.

696 And for those who know Hoon, `[[4 0 1] 51]` should al-
697 ready be looking a lot like `[battery payload]...` that's not
698 a coincidence. We're starting to see the first glimpses of how
699 Hoon's cores, arms and subjects/subject mutations flow out of
700 the fundamental structure of Nock.

701 We also see how Hoon/Nock lend themselves well to throw-
702 ing around chunks of code, and adjusting the subject as neces-
703 sary to create the correct subject/environment against which
704 to run that code.

705 2.9 Summary of Fundamental Opcodes

706 You now have seen all the fundamental functions/opcodes in
707 Nock. In the next part, I'll introduce the remaining functions,
708 which no longer need pseudocode: we have enough scaffold-
709 ing now to build the rest of Nock from Nock itself. Instead,
710 these new opcodes will just be shortcuts/macros/code expan-
711 sions building on opcodes 0-5.

712 We'll also start to connect Nock to Hoon, and see how most
713 of the fundamental (and slightly weird) features of Hoon flow
714 directly from Nock's structure, and make a lot more sense in
715 combination with it.

716 3 Composite Opcodes

717 A word of encouragement: we’re done with the hard part now.
718 Every Nock function we learn in this section will be built from
719 pieces in the preceding section.

720 None of these new functions are *necessary* to make Nock
721 work. All of them except Nock 10 and 11 are syntactic sugar
722 that is part of the Nock definition, and must be built into every
723 correct Nock implementation. If you have seen macros or code
724 expansions in other languages, that’s another word for what’s
725 happening here.

726 Nock 10 makes it easier to replace a memory value some-
727 where in a tree, and Nock 11 allows passing hints to the inter-
728 preter.

729 One point of clarity: this syntactic sugar/code-expansion
730 system is *not* extensible. This means that you can’t invent your
731 own Nock opcodes and still have that language be Nock; you’re
732 making a higher-level language on top of Nock at that point.

733 In fact, that’s not a bad way to think of Hoon: it’s a higher-
734 level language that adds missing syntactic sugar/macros and
735 human-readable names to Nock. (That’s not the whole story,
736 but it’s a decent mental peg to initially hang Hoon on.)

737 Nock is intentionally very, very small, such that you can
738 always walk through and analyze what is happening in a block
739 of code if you know Nock’s opcodes.

740 3.1 An Aside About the 1 (Quoter Function)

741 You’re going to see the 1 opcode (the “Quoter”) appear in a lot
742 of examples below for a simple reason: the 6-9 opcodes expect
743 formulas in a lot of places, not just atoms. And, as we’ve seen
744 already, formulas have to be cells.

745 Whenever a formula is required, but you really just want
746 to return a number, you use the quoter function.

747 3.1.1 Example: Incrementing a Constant

748 :: We just want to run 4 on the number 5,
749

```

750 :: but 4 expects a formula after it, so we use [1 5]
751 > .* (0 [4 1 5])
752 6
753

```

754 3.1.2 Example: Comparing a Memory Slot to a Constant Value

```

755
756 :: we grab memory slot 2
757 :: then it has to be compared to the result of a
758     formula
759 :: so we just use the formula [1 23] to return 23
760 > .* ([23 45] [5 [0 2] [1 23]])
761 0
762

```

763 3.2 6, "If/Else" Conditional Branch

764 The remaining Nock opcodes are “sugar”. This means that the
765 functions in this next part will use only `*` and Nock code in
766 their pseudocode. For example, here is the code for 6, the
767 “If/Else” function.

```

768
769 * [a 6 b c d]      * [a * [[c d] 0 * [[2 3] 0 * [a 4 4 b]]]]
770

```

771 The prose explanation of what’s happening is straightforward:
772

- 773 1. Evaluate `b` against the subject `a` (`* [a b]`) to see whether
774 it’s `0/true` or `1/false`.
- 775 2. If `b` equals `0/true`, run formula `c` against subject `a`.
- 776 3. If `b` equals `1/false`, run formula `d` against subject `a`.
- 777 4. If `b` is not equal to `0` or `1`, the code crashes, for reasons
778 we’ll see in the code explanation.

779 3.2.1 Code Explanation (Way More Fun)

780 OK, so that’s the English version. The code explanation (the
781 right side of the definition) is *really fun* now that we know the
782 basic Nock opcodes.

783 The pseudocode has four nested “[subject formula]’s, so
784 I’m going to unwrap those to the bottom, and then build it up
785 again. The layers are, in order:

- 786 1. `*[a ...]`, i.e. subject `a` evaluated against that long for-
787 mula starting with `*[c d] @ ...]`.
- 788 2. `*[c d]`, i.e. subject `c` evaluated against the formula start-
789 ing with `[2 3 ...]`.
- 790 3. `*[2 3]`, i.e. subject `2` evaluated against the lowest-level
791 formula.
- 792 4. Finally, subject `a` evaluated against formula `[4 4 b]`.

793 3.2.2 Step 4

794 Remember, our English explanation was “see if `*[a b]` is true
795 or false, and do different actions depending on that. Step Four
796 is that check.

797 Let’s say we have `a` as `59`, and `b` as the quoted value `@/true`.

```
798 *[a 4 4 b]
799 :: a: 59
800 :: b: [1 0]
801 *[59 4 4 [1 0]]
802 :: substitute out the two increment operators
803 ++*[59 [1 0]]
804 :: ignore subject, return quoted value 0
805 ++(0)
806 2
```

807 In summary, because `*[a b]` evaluated to `@/true`, we get
808 the number `2`. If `*[a b]` had been `1/false`, we’d get `3` (because
809 we’d evaluate `++(1)`).

810 What the heck? Why are we getting `2` or `3` back? How
811 does that help us?? Well, Nock uses that returned `2` or `3` as a
812 *memory slot*, and executes the code in that memory slot.

813 3.2.3 Step 3

814 Now we go up to step 3, with the subject `[2 3]` evaluated
815 against our return value from step 4. Let’s imagine that `2` had

816 been returned:

```
817 :: remember, "result-of-step-4" was 2
818 *[[2 3] @ result-of-step-4]
819 *[[2 3] @ 2]
820 :: get memory slot 2
821 2
```

822 This grabs memory slot 2 if $[a \ b]$ was true, and memory
823 slot 3 if $[a \ b]$ was false.

824 Wait, isn't this redundant? We are just using our 2 or 3
825 generated in step 4 to generate a 2 or a 3. Seems dumb.

826 The answer is that we make sure that a crash happens if
827 $[a \ b]$ yields any answer other than 0/true or 1/false. If $[a \ b]$
828 returned 10, for example, we'd have the following code in
829 step 3:

```
830 *[[2 3] @ 10]
831 :: there's no slot 10
832 CRASH!!
```

833 This is exactly what we want: the program crashes unless
834 we are doing a loobean⁵ test that returns a 0 or 1 (converted to
835 an indicial 2 or 3) in step 4.

836 3.24 Step 2

837 We now have our validated 2 or 3 to plug into step 2. Let's
838 imagine c is the simple formula $[0 \ 1]$ and d is $[1 \ 203]$.

```
839 :: if step 3 returned "2" (true)
840 *[[[0 1] [1 203]] @ 2]
841 [0 1]
```

842 Not much to see here: we just grab memory slot 2 or 3
843 depending on whether our initial b was true or false.

844 3.25 Step 1

845 And now we're back at the top level, where we just use whichever
846 formula we yonked in step 2 and run it against a .

⁵A boolean truth value but with the typical truth values of 0 and 1 reversed to 0/true and 1/false.

```

847 *[a formula-from-step-2]
848 :: let's say we returned formula [0 1]
849 :: our original a, from step 4, was 59
850 *[59 [0 1]]
851 59

```

852 3.2.6 Example: Code Expansion of 6

```

853
854 > .*(1 [6 [0 1] [0 1] [4 0 1]])
855 :: PSEUDOCODE
856 :: *[a *[[c d] 0 *[[2 3] 0 *[a 4 4 b]]]]
857 *[1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 *[1 4 4 [0 1]]]]]
858 :: factor out the 4 opcodes
859 *[1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 **[1 [0 1]]]]]
860 :: b evaluates to 1 (yank memory slot 1)
861 *[1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 **([1])]]]
862 :: evaluate the two increments
863 *[1 *[[[0 1] [4 0 1]] 0 *[[2 3] 0 3]]]
864 :: get memory slot 3 from [2 3]
865 *[1 *[[[0 1] [4 0 1]] 0 3]]
866 :: [0 3] means get memory slot 3 from the subject
867      (formula [4 0 1])
868 *[1 [4 0 1]]
869 :: factor out the 4 opcode
870 **[1 [0 1]]
871 +(1)
872 2
873

```

874 3.2.7 Summary of 6

875 Now I'm going to make you a little sad. Most Nock interpreters
876 don't do this whole awesome code expansion. They just see 6,
877 and implement an if/else check with a crash if *[a b] isn't a
878 loobean.

879 However, the pattern of storing chunks of code in memory
880 and pulling them out when you want them is the most impor-
881 tant part of Nock. In fact, those of you who know Hoon are
882 probably already seeing the makings of the “core” code pat-
883 tern.

3.3 7, the "Composition" Opcode

7 is so simple we barely need to spend time on it. All it does is create a new subject/environment (using 2), and immediately runs a formula against that new subject. Let's compare it to 2:

```

:: opcode 7
*[a 7 b c]  *[*[a b] c]

:: opcode 2
*[a 2 b c]  *[*[a b] *[a c]]

```

This is almost exactly the same except with c instead of *[a c]. This means that you can write a "subject-changing" formula in b, and then run a simple function against that in c.

3.3.1 Example: Opcode 2 vs. Opcode 7

```

:: using opcode 2
> .*([23 45] [2 [0 3] [1 4 0 1]])
46

:: using opcode 7 -- we get to remove the quoter
  opcode "1"
:: wow amazing /sarcasm
> .*([23 45] [7 [0 3] [4 0 1]])
46

```

Opcode 7 is clearly trivial, so why does it exist? It allows clean expression of function composition: this is really $c(b(a))$, where a (the subject) is the initial argument, and b and c are functions. This pattern comes up a ton, and it's nice to not use 1s everywhere.

3.4 8, the "Variable Push" Opcode

If you're ever writing Nock and think "how can I add a new variable to the subject?", opcode 8 is what you want. The new variable can be based on either the existing subject, or be a new value you add on.

```

*[a 8 b c]          *[[*[a b] a] c]

```

923 This says to run `*[a b]`, and then make that the head of a
 924 new subject, with the old subject `a` as the new subject's tail.

925 3.4.1 Example: Add Variable as a Copied Value from an Existing Sub- 926 ject

927 In English, the below code first yanks the variable from mem-
 928 ory slot 3, copies it to the head of a new subject, and then in-
 929 crements the value in that new head.

```

930 > .*([67 39] [8 [0 3] [4 0 2]])
931 40
932
933 :: PSEUDOCODE
934 *[[*([67 39] [0 3]) [67 39]] [4 0 2]]
935 :: yanks mem slot 3 and pins it to the front of the
936        old subject
937 *[[39 [67 39]] [4 0 2]]
938 **[[39 [67 39]] [0 2]]
939 +(39)
940 40
  
```

942 3.4.2 Example: Add Variable as a New Value

```

943 > .*([67 39] [8 [1 0] [4 0 2]])
944 1
945
946 :: PSEUDOCODE
947 *[[*([67 39] [1 0]) [67 39]] [4 0 2]]
948 *[[[0 [67 39]] [4 0 2]]
949 **[[[0 [67 39]] [0 2]]
950 +(0)
951 1
  
```

953 Why do we want this? In the first example, we pin a copy
 954 of a value so that we can manipulate it without changing the
 955 original. In the second example, we add a `0` to the front; maybe
 956 we want to increment it until some condition is met?

957 The above should be starting to feel *very* Hoon-ish: we're
 958 a minor code transform away from pinning new values to the
 959 head of a payload.

960 3.5 9, Run a Stored Procedure Arm in a Core

961 We're almost at full Hoon now, although still at a very low/raw
962 level. 9 looks a little complicated...

964 $\star[a \ 9 \ b \ c] \ \star[\star[a \ c] \ 2 \ [0 \ 1] \ 0 \ b]$

966 ... but in English, this is just saying:

- 967 1. Use formula c to make a new subject from a ($\star[a \ c]$).
- 968 2. Grab the formula located at memory slot b in that new
969 subject.
- 970 3. Run that formula against the new subject $\star[a \ c]$.

971 The fact that the above description has the words “make a new
972 subject” tells us right away that it’s syntactic sugar for opcode
973 2, and that’s indeed what we see in the pseudocode.

974 3.5.1 Example: Using 9 to Run an Increment Arm

975 The initial expression here likely looks cryptic, but if you fol-
976 low the pseudocode, it will become clear.

977 To stay oriented, remember that, if we’re thinking of 9 as
978 $\star[a \ 9 \ b \ c]$, then

979 a : 45
980 b : 2
981 c : $[[1 \ 4 \ 0 \ 3] \ 0 \ 1]$

```

982 > .*(45 [9 2 [1 4 0 3] 0 1])
983 46
984 :: PSEUDOCODE
985  $\star[45 [9 \ 2 \ [1 \ 4 \ 0 \ 3] \ [0 \ 1]]]$ 
986  $\star[\star[45 [1 \ 4 \ 0 \ 3] \ [0 \ 1]] \ 2 \ [0 \ 1] \ 0 \ 2]$ 
987 :: new subject is  $[[4 \ 0 \ 3] \ 45]$ 
988 :: that is a formula to increment mem slot 3 in the
989 :: head; 45 in the tail
990  $\star[[[4 \ 0 \ 3] \ 45] \ 2 \ [0 \ 1] \ 0 \ 2]$ 
991 :: now we expand opcode 2
992  $\star[\star[[[4 \ 0 \ 3] \ 45] \ 0 \ 1]$ 
993  $\star[[[4 \ 0 \ 3] \ 45] \ 0 \ 2]]$ 
```

```

995 :: mem slot 2 of the subject becomes the new formula
996 *[[[4 0 3] 45] 4 0 3]
997 **[[[4 0 3] 45] 0 3]
998 :: grab mem slot 3
999 +(45)
1000 46
1001

```

1002 We start above with a subject that is not a core; it's just the
1003 atom 45. The code for `c` is then:

```
1004 [[1 4 0 3] [0 1]]
```

1005 This formula uses the Cell Maker (Distribution Rule) to insert
1006 [4 0 3] as the head of the new subject, and puts mem
1007 slot 1 of the old subject as the tail, so we get a new subject of
1008 [[4 0 3] 45].

1009 What we do next is use `b` (here, 2) to select memory slot 2
1010 from that new subject. Memory slot 2 is a formula: [4 0 3].
1011 We run that formula against the new subject.

1012 3.5.2 Key Point/Possible Confusion

1013 When I first saw 9, I thought: “why didn’t they just use 2 to
1014 make the transformed subject? Why is there an extra step to
1015 use [0 1] to pull the new `*[a c]` subject? Why can’t we do
1016 `*[a 2 c [0 b]]`?

1017 The answer is that we actually use the `*[a c]` subject *twice*:

- 1018 1. We extract from it the formula located at memory slot `b`.
- 1019 2. Then we run *that* formula against the `*[a c]` subject.

1020 If we just did `*[a 2 c [0 b]]`, then [0 b] would try to
1021 look up the `b` memory slot in `a`, NOT `*[a c]`. So 9 gives us
1022 one extra step to set up the core itself.

1023 3.5.3 How to Think of 9

1024 I like to think of 9 as having three parts:

- 1025 1. `c`: our formula to set up the subject, making a new sub-
1026 ject.

- 1027 2. b: the memory slot in our new subject where an arm is.
 1028 3. Run the arm located at b against the new subject.
 1029 You can think of this as setting up a subject, pulling a stored
 1030 procedure from it, and then running that procedure against the
 1031 subject.

1032 3.6 10, Replace a Memory Slot

1033 Before explaining 10, we need to introduce a new operator, #.
 1034 # is the “edit” operator. It has the form
 1035

 1036 #[mem-slot new-val target-tree]
 1037

1038 It replaces the memory slot mem-slot in target-tree with
 1039 new-val.

1040 :: Example
 1041 #[2 [4 5] [99 88 77]]
 1042 [[4 5] 88 77]

1043 The pseudocode for 10 is:

1044

 1045 * [a 10 [b c] d] # [b * [a c] * [a d]]
 1046

1047 In English, this first calculates * [a c] and * [a d], and
 1048 then replaces memory slot b in the latter with the result of the
 1049 former.

1050 3.6.1 Example: Using 10 to Replace a Memory Slot

1051

 1052 > .*(50 [10 [2 [0 1]] [1 8 9 10]))
 1053 [50 9 10]
 1054 :: PSEUDOCODE
 1055 *[50 [10 [2 [0 1]] [1 8 9 10]]]
 1056 #[2 *[50 0 1] *[50 1 8 9 10]]
 1057 #[2 /[1 50] [8 9 10]]
 1058 #[2 50 [8 9 10]]
 1059 :: expanded as far as we can, now do the edit
 1060 :: replace mem slot 2 of [8 9 10] with the value 50
 1061 [50 9 10]
 1062

1063 We will defer discussion of opcode 11 to a later section be-
 1064 cause it is not part of strict Nock semantics.

1065 3.6.2 Towards Hoon

1066 If “grab a chunk of code from a subject and then run it against
 1067 the subject” sounds a lot like getting an arm from a core in
 1068 Hoon, that’s because it is. Nock opcode 9 is hand-in-glove
 1069 with Hoon’s core/arm structure. The “new subjects” created by
 1070 `*[a c]` are cores, and the things selected from memory slots
 1071 by `b` are arms.

1072 3.7 Real Nock Code

1073 To finish this off and take your new powers for a spin, let’s look
 1074 at some real Nock code from the wild. I got the below example
 1075 from the Dojo. It’s a mold: a function that takes a noun and
 1076 returns it if it’s the correct type, and crashes if not. The mold
 1077 here checks whether the input noun is a loobean (`0/true` or
 1078 `1/false`).

1079 3.7.1 Example: A Loobean Mold

```

1080 :: loobean mold gate
1081 > =loobean-mold ?
1082
1083
1084 :: below output truncated for our purposes here
1085 > loobean-mold
1086 < 1.fxf ... >
1087
1088 :: grab the code for the battery
1089 > -.loobean-mold
1090 [8 [6 [5 [1 0] 0 6] [1 0] 6 [5 [1 1] 0 6] [1 1] 0 0]
1091 8 [5 [0 14] 0 2] 0 6]
1092
```

1093 So we:

- 1094 1. Assign the mold gate denoted by `?` to the face `loobean-mold`.
- 1095 2. Examine what’s inside, seeing that the head is an arm.
- 1096 3. Print out that source code by calling the head.

1097 The below code acts as a formula which evaluates its sample,
 1098 and returns it if it's a loobean or crashes otherwise. Let's see
 1099 how that works.

1100 `[6 [5 [1 0] 0 6] [1 0] 6 [5 [1 1] 0 6] [1 1] 0 0]`

1101 We start with `6`, which means this is an if-else. The true/false
 1102 test is the next element:

1103 `[5 [1 0] 0 6]`

1104 This compares the quoted value `0` (from `[1 0]`) with the value
 1105 at memory slot `6` in the subject (`[0 6]`).

1106 What is the subject and what is at mem slot `6`? This code is
 1107 the arm of a gate, so the subject is that gate/core: `[battery`
 1108 `sample payload]`. We are looking at the battery right now,
 1109 and so `[0 6]` yanks the head of the tail, or the `sample`!

1110 We know the sample will be a noun that we're testing for
 1111 loobeanness, so this code starts by seeing whether the sample
 1112 is the value `0`/true. If it is, the next element is `[1 0]`, so we
 1113 return `0` if the sample is `0`.

1114 Otherwise, we run the second branch of the `if-else` con-
 1115 ditional:

1116 `[6 [5 [1 1] 0 6] [1 1] 0 0]`

1117 This is also an opcode `6` if-else. Once again, it compares
 1118 something to the value at mem slot `6` (the sample), but this
 1119 time it checks whether that is the value `1`. If it is, it runs the
 1120 formula `[1 1]` to return `1`.

1121 If not, it means our input was neither `0` nor `1` and is not a
 1122 loobean, so we run the formula `[0 0]`, which always crashes
 1123 (as there is no memory slot `0`).⁶ This is exactly what we want—
 1124 crash if the sample is not a loobean!

1125 The second part of the formula which is composed with an
 1126 `8` is:

1127 `[8 [5 [0 14] 0 2] 0 6]`

1128 This checks whether the battery of the current gate is equiv-
 1129 alent to the head of the context; if so, it returns the sample. This

⁶Recall that in the binary tree addressing scheme of Nock, `1` refers to the entire noun, `2` to its head if it exists, and `3` to its tail if it exists.

1130 is an artifact of a type promotion in Hoon which converts a bare
1131 loobean value to a constant by pinning a constant as an exam-
1132 ple first (i.e., at slot 14). (It's completely extraneous in Nock
1133 itself, however.)

1134 3.8 Summary

1135 In this section, we've seen how to use almost all of the remain-
1136 ing opcodes, which build Nock up to a slightly more expres-
1137 sive level. We also saw how Hoon cores start to arise pretty
1138 naturally out of the Nock primitives, especially 8 and 9. Then
1139 we walked through real production Nock code to show that
1140 everything we've learned so far works exactly as expected in
1141 the wild. In the next section, you'll learn how to write real
1142 programs in Nock, and compose those programs to make new
1143 ones. Finally, hopefully you now see that, whatever its other
1144 limitations, Nock is not particularly obscurantist, and is fairly
1145 straightforward to parse, once you understand its syntax and
1146 idioms.

1147 4 Interlude

1148 In this section, we cover some loose ends deriving from the
1149 previous discussion.

1150 4.1 Order of Operations

1151 At this point, we can answer the question of what order Nock
1152 evaluates expressions in. The answer is straightforward: it
1153 starts with code that has only one pseudocode operator and
1154 all Nock code, as below:

```
1155 *[50 4 4 [5 [0 1] [1 50]]]
```

1156 The interpreter can then be thought of as moving left-to-right,
1157 expanding according to its rules, until it can't expand any fur-
1158 ther into pseudocode:

```
1159 *[50 4 4 [5 [0 1] [1 50]]]
```

```
1160 **[50 4 [5 [0 1] [1 50]]]
```

```

1161 ++*[50 5 [0 1] [1 50]]
1162 ++=[*[50 [0 1]] *[50 [1 50]]]
1163 ++=[/[1 50] 50]

```

1164 Once we get to that last line, ++=[/[1 50] *[50 [1 50]]],
 1165 no further expansion into pseudocode is possible.

1166 At this point, the interpreter expands “inside-out”, starting
 1167 from the deepest operators. In this case, those are /[1 50]
 1168 and 50, so it does those:

```

1169 ++=[50 50]

```

1170 Then it moves to the next-furthest-inside operator: =, and then
 1171 to the + operators:

```

1172 ++=[50 50]
1173 ++(0)
1174 +(1)
1175 2

```

1176 4.1.1 Summary of Nock Order of Operations

- 1177 1. On the first pass, when we have Nock code `*(nock-code)`,
 1178 we expand left-to-right into pseudocode.
- 1179 2. On the second pass, when we’ve fully expanded pseu-
 1180 docode, we evaluate operators from the inside-out.

1181 4.2 11, Hints and Side Effects for the Interpreter

1182 Opcode 11 lets you compute side effects and pass informa-
 1183 tion to Nock’s interpreter, for that interpreter to do what it
 1184 wants. The most common uses are things like jets and debug-
 1185 ging prints. In theory, this opcode is fairly dangerous, since it
 1186 tells the interpreter that it’s free to do anything.

1187 There are two different versions of 11, depending on whether
 1188 the head following 11 is an atom or a cell

```

1189 :: head is an atom (b)
1190 *[a 11 b c] *[a c]
1191
1192 :: head is a cell ([b c])
1193 *[a 11 [b c] d] *[[*[a c] *[a d]] 0 3]

```

- 1194 This pseudocode lets 11 handle two separate use cases:
- 1195 1. We just want the interpreter to process a message in its
 - 1196 own language/environment.
 - 1197 2. We want the interpreter to compute some Nock code
 - 1198 with Nock semantics before processing the hint.

1199 4.2.1 Option #1: Interpreter Processes a Message

1200 Imagine the following Nock:

```
1201 . *([50 51] [11 369 0 2])
```

1202 This passes the value 369 to the interpreter, where maybe it
 1203 would be the key in a hashmap. The value associated with the
 1204 could be the function to debug print the whole operation. Or
 1205 the value could be a jet function that specifically knows how
 1206 to compute the 0 opcode faster when it's followed by a 2.

1207 After the interpreter does whatever, it needs to then discard
 1208 the value 292.989, and computes

```
1209 > . *([50 51] [0 2])
1210 50
```

1211 4.2.2 Option #2: Run Nock Code and then Process a Message

1212 In this case, instead of passing a static value to the interpreter,
 1213 we pass it a Nock computation against the subject. The result
 1214 of that computation will then be available to the interpreter to
 1215 use as a message.

1216 4.2.3 Practical Use of 11: Jet Registration

1217 Jet registration in Hoon usually looks like this:

```
1218 ++ my-function
1219 ~ / %my-function
1220 | = var = *
1221 ...
1222
```

1224 ~ / “fassig” uses Nock 11 to pass the value %my-function to the
 1225 interpreter, where it is associated with the C code that produces
 1226 the product of my-function.

4.3 Nock in Hoon

In Hoon, you can acquire the Nock code for any function by using the `!=` “zaptis” rune. Since this produces Nock code complete with gate calls, you may also find it helpful to cast a data type as a noun using `^ - *`, which will show you the raw Nock code for that value.

As you have seen, you can also evaluate raw Nock using the `. *` “dottar” rune, which evaluates to Nock opcode 2. Other Nock runes in Hoon include `. +`, which is opcode 4; `. =`, which is opcode 5; and `. ?`, which is opcode 3. Fake opcode 12 is not actually Nock, but is used in a virtualization context to resolve values that may not be available in the current subject.

5 The Core as Design Pattern

At this point, you now know everything there is to know about Nock syntax and basic code manipulation. However, if I asked you to go write a simple program in Nock, you’d probably have a hard time getting started. It’s really helpful to see some examples of how Nock programs are designed and patterns that tend to recur in them.

We can set variables to formulas in the Urbit Dojo. We will use that in this lesson to make code easy to follow, and to show how Nock programs are made out of discrete chunks.

```
> =increment-formula [4 @ 1]
> .*(50 increment-formula)
51
> =increment-formula
```

5.1 Building a Core Manually

Every Nock program is just a formula that takes in a subject (argument). In order to run full programs against a subject, we follow the following pattern:

1. First pass: set up the dominos inside an opcode like 2 (or 7/8/9). This will create a core, and evaluate it on the second pass.

1260 2. Second pass: evaluate that core to get our result.

1261 The pseudocode for opcode 2, this looks like:

```
1262        *[subject 2 core-formula call-formula]
```

1265 Almost always, however, we'll use the 9 opcode (which is
 1266 sugar on top of 2) to actually start our programs running, since
 1267 it supports the concept of “cores” and “arms” very naturally.
 1268 (For those who know Hoon, this “set up a core and run it” pat-
 1269 tern should feel similar to | ^ or =>.)

1270 5.1.1 Example: A Simple Increment Gate

1271 Let's take code to increment the subject and turn it into a Hoon-
 1272 style gate (a core with one arm and a sample). The code itself
 1273 is simple: [4 0 1].

1274 Our code expects the subject to be an atom, but Hoon gates
 1275 expect their subject to be the gate itself, which has the form
 1276 [battery payload]. So we need to adjust our formula to ac-
 1277 cept this subject format.

```
1278        > =inc [7 [0 3] [4 0 1]]
1279        :: verify that it increments value in the subject's
1280        tail (payload)
1281        > .*([1 74] inc)
1282        75
```

1285 We put our formula in the head, and put a default (“bunted”)
 1286 sample of 0 as the payload.

```
1287        > =inc-gate [inc 0]
1288        > inc-gate
1289        [[7 [0 3] 4 0 1] 0]
```

1290 Arms are best invoked using the 9 opcode. Here we use the
 1291 “standard” version of it: [9 2 0 1]. Recall that 9 is $\star[a \ u \ b$
 1292 $c]$, where c is the “new subject” formula, and b is the memory
 1293 slot of that new subject to take an arm from. Here we use [0
 1294 1] to keep our subject (the gate) unchanged.

```
1295        > .*(inc-gate [9 2 0 1])
1296        1
```

```

1297     To “pass a parameter” to our gate, we use the 10 opcode to
1298     edit the subject (the inc-gate core) and replace the payload
1299     with a value. We make our subject [inc-gate val-to-increment]:

1300     :: notice how opcode 10 can replace the last element
1301         of inc-gate
1302     :: the tail of inc-gate was 0; now it's 36
1303     > .*([inc-gate 36] [10 [3 [0 3]] 0 2])
1304     [[7 [0 3] 4 0 1] 36]
1305
1306     :: now use that to set up the subject with a new
1307         tail (sample)
1308     :: and call arm in memory slot 2 (our increment
1309         formula)
1310     > .*([inc-gate 36] [9 2 10 [3 [0 3]] 0 2])
1311     37
1312     > .*([inc-gate 562] [9 2 10 [3 [0 3]] 0 2])
1313     563

```

1314 5.1.2 Summary

1315 What was the point of this? We took a perfectly good incre-
 1316 ment formula that worked on atoms, and made it more com-
 1317 plicated for the same result. *Lame!*

1318 In fact, however, this setup will come in really handy when
 1319 we have multiple arms in our core, not just one. We’ve made
 1320 it easy to set up a clean environment for increment to find its
 1321 sample in, no matter what other data is present. This will come
 1322 in *very* handy in our third example program.

1323 First, however, let’s look at putting a more complicated for-
 1324 mula inside a gate. (Hint: the process is the same).

1325 5.1.3 Example: A Decrement Gate

1326 Here we take a piece of code that decrements the subject when
 1327 the subject is an atom, and turn it into a gate. For now, we
 1328 will take it as a given that the decrement function given below
 1329 works. (The code will crash when the input is a cell or `0`)

```

1330 :: the code: set a Dojo face =dec
1331

```



```

1332 > =dec [6 [5 [0 1] [1 0]] [0 0] [6 [3 0 1] [0 0] [8
1333   [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0
1334   6] 0 7] 9 2 0 1]]]
1335
1336 :: testing in the Dojo
1337 > .*(100 dec)
1338 99
1339
1340 :: with non-atom input we crash
1341 .*([100 101] dec)
1342 1343 dojo: hoon expression failed

```

1344 Let's re-jigger the decrement formula to accept the subject
 1345 in format [battery payload]:

```

1346 :: gets value in mem slot 3 and applies the
1347   decrement formula to it
1348 > =dec-arm [7 [0 3] dec]
1349
1350 :: verify that dec-arm decrements the value in the
1351   tail
1352 > .*([1 88] dec-arm)
1353 87
1354
1355 :: output the full dec-arm formula code
1356 > dec-arm
1357 [7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1358   0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1359   0 7] 9 2 0 1]

```

1360 Next, assemble the core:

```

1361 :: all we need to do is add the sample! (we give it a
1362   default value of 0)
1363 > =dec-gate [dec-arm 0]
1364 > dec-gate
1365 [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1366   0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1367   0 7] 9 2 0 1] 0]
1368 [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1369   0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1370   0 7] 9 2 0 1] 0]

```

```

1371     As in the first example, we use the 10 opcode to edit the
1372     subject (the dec-gate core) and replace the payload with a
1373     value. We make our subject [dec-gate val-to-decrement].

1374     :: confirm that our 10 code places 36 in the tail
1375     > .*([dec-gate 36] [10 [3 [0 3]] 0 2])
1376     [[7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0 0] 8 [1
1377         0 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2] [4 0 6]
1378         0 7] 9 2 0 1] 36]
1379
1380     :: now use that to set up the subject with a new
1381         tail (sample)
1382     :: and call arm in memory slot 2 (our decrement
1383         formula)
1384     > .*([dec-gate 36] [9 2 10 [3 [0 3]] 0 2])
1385     35
1386     > .*([dec-gate 562] [9 2 10 [3 [0 3]] 0 2])
1387     561
    
```

1388 5.1.4 Example: Comparing Two Numbers

1389 In this example we will build a Nock core that compares two
 1390 numbers and returns whether they are greater than, less than,
 1391 or equal to each other.

1392 The program specification is to take two numbers `a` and `b`
 1393 and return:

- 1394 • 0 if `a == b`.
- 1395 • 1 if `a > b`.
- 1396 • 2 if `a < b`.

1397 The trick is that the only mathematical operators we have
 1398 for this are the 4 opcode for increment, 5 for equals, and a
 1399 decrement Nock function that we will supply. The basic algo-
 1400 rithm is:

- 1401 • If `a == b`, return 0.
- 1402 • If `a == 0`, return 2. (I.e, `a < b`).
- 1403 • If `b == 0`, return 1. (I.e., `a > b`).

- Recur with a and b both decremented.

The above works because if a is smaller than b, it hits 0 first. If b is smaller than a, it hits 0 first. If they are equal, we never decrement in the first place, so there are no issues with numbers going negative.

We will construct a core that has two arms: one arm with the algorithm logic, and a second arm with the decrement gate from the second example. These arms will live in the head of the core (think “battery”).

We will put our variables a and b in the tail of the core (think “payload”). They will be updated before the core is called each time. (This is similar to a trap or door in Hoon).

The final core we will create will have the structure [battery payload], which can be broken down as:

```
[[main-logic dec-gate] a b]
```

Once the core is set up, we’ll invoke the algorithm logic arm to set it running. We will build the main logic “inside-out”, by first defining the “recur” code, and then inserting it into our if/else tests that implement the algorithm.

This recursion logic assumes the decrement gate lives in the tail of the battery (the battery is the head), i.e. [0 5]. We are using [10 [3 [0 memory-slot]] 0 5] to extract the decrement gate from our core, and then we invoke it with [9 2 ...].

This is identical to what we did in the second example—the only difference is that we now do it twice: once for memory slot 6 (a) and again for memory slot 7 (b).

```
> =dec-a [9 2 10 [3 [0 6]] 0 5]
> =dec-b [9 2 10 [3 [0 7]] 0 5]

:: we can test our formulas by making a
:: dummy core (a = 33, b = 77)
:: first formula in battery just returns
:: the current subject
> =dummy-core [[[0 1] dec-gate] 33 77]
> .* (dummy-core dec-a)
32
```

```

1441 > .*(dummy-core dec-b)
1442 76

```

The recursion formula here operates by invoking its subject—
itself—and applying the `dec-a` and `dec-b` formulas to it. Then
we use opcode 9 for the new subject's setup: replace payload
with new `a` and `b` run the current core

```

1447 > =recur [9 4 [[0 2] dec-a dec-b]]
1448
1449 :: test it -- we want to see same battery,
1450 :: with payload being replaced with 32 and 76
1451 > .*(dummy-core recur)
1452 [[[[0 1] [7 [0 3] 6 [5 [0 1] 1 0] [0 0] 6 [3 0 1] [0
1453         0] 8 [1 0] 8 [1 6 [5 [0 7] 4 0 6] [0 6] 9 2 [0 2]
1454         [4 0 6] 0 7] 9 2 0 1] 0] 32 76]
1455
1456 :: clean up the dummy-core
1457 > =dummy-core

```

The main logic is then:

```

1458
1459 :: check whether a is 0; return 2 if true
1460 :: else check whether b is 0; return 1 if true
1461 :: else recur
1462 > =main-logic [6 [5 [0 6] [1 0]] [1 2] [6 [5 [0 7] [1
1463         0]] [1 1] recur]]
1464
1465
1466 :: test (we aren't doing the equality case here)
1467 :: a < b
1468 > .*([[main-logic dec-gate] 9 10] [9 4 0 1])
1469 2
1470 :: a > b
1471 > .*([[main-logic dec-gate] 10 9] [9 4 0 1])
1472 1
1473

```

We build an outer test for whether `a==b`, and then if it's
not, we use the 1 opcode to return the core.

```

1474 > =battery [main-logic dec-gate]
1475
1476 :: payload for our core is memory slots 2 and 3 (a
1477 and b)
1478 > =comparison [6 [5 [0 2] [0 3]] [1 0] [9 4 [[1
1479         battery] [0 2] 0 3]]]
1480
1481

```

```

1483
1483 :: test our 3 cases
1484 > .*([0 8] comparison)
1485 2
1486 > .*([0 0] comparison)
1487 0
1488 > .*([8 0] comparison)
1489 1

```

1490 We can take this code and use it *anywhere* that our subject
 1491 is a cell of two numbers, and it will “just work”.

```

1492 [6 [5 [0 2] 0 3] [1 0] 9 4 [1 [6 [5 [0 6] 1 0] [1 2]
1493 6 [5 [0 7] 1 0] [1 1] 9 4 [0 2] [9 2 10 [3 0 6] 0
1494 5] 9 2 10 [3 0 7] 0 5] [7 [0 3] 6 [5 [0 1] 1 0]
1495 [0 0] 6 [3 0 1] [0 0] 8 [1 0] 8 [1 6 [5 [0 7] 4 0
1496 6] [0 6] 9 2 [0 2] [4 0 6] 0 7] 9 2 0 1] 0] [0 2]
1497 0 3]

```

1498 One common convention used in languages like Hoon is
 1499 to denote Nock rules as constants, e.g. %0 (corresponding to
 1500 our 0). This makes it much more straightforward to interpret
 1501 hand-rolled Nock code with addresses and constants present.

1502 6 Conclusion

1503 After this tutorial, you should be prepared to interpret Nock
 1504 code when it is produced by Hoon. Although there is only
 1505 rarely any call to write Nock code directly, you should now
 1506 have a good understanding of how to read and interpret it. As
 1507 you write your own interpreters in the future, you may even
 1508 have call to hand-roll Nock expressions on occasion.